

CSC424 System Administration

Instructor: Dr. Hao Wu

Week 6 Bash Scripting - 1

Outline

- Filename Matching
- Regular Expression
- Linux Text Processing Tools
- Variables
- Arithmetic
- Control Flow
- Loops
- Function

Filename Matching

- Wildcards (meta characters) are symbols or special characters that represent other characters. You can use them with any command such as ls, cp, mv, rm to perform file operations matching a given criteria.
- Three main wildcards in Linux:
 - * : matches one or more occurrences of any character, including no character
 - ? : represents or matches a single occurrence of any character
 - [] : matches any occurrence of character enclosed in the square bracket. It is possible to use different types of characters (alphanumeric characters): numbers, letters, other special characters etc.)

Zero or multiple characters matching (*)

1. If we want to list all the files start with 'a':

`ls a*`

```
[hao@localhost sbin]$ ls a*
accessdb      agetty      arpd         auditd      authconfig
addgnupghome  alternatives arping       augenrules  authconfig-tui
addpart       anacron     audispd      aureport    autrace
adduser       applygnupgdefaults auditctl     ausearch    avcstat
```

2. If we want to list all the files end with 'a'

`ls *a`

```
[hao@localhost sbin]$ ls *a
thin_delta  xfs_quota
```

Zero or multiple characters matching

1. If we want to list all the files start with 'a' and end with 't':
`ls a*t`

```
[hao@localhost sbin]$ ls a*t  
addpart  aureport  avcstat
```

Single character matching (?)

1. If we want to list a file whose name contains 2 characters and starts with 's'

ls s?

```
[hao@localhost sbin]$ ls s?  
ss
```

2. If we want to list a file whose name contains 2 characters and ends with 's'

ls ?s

```
[hao@localhost sbin]$ ls ?s  
ss
```

Single character matching (?)

1. If we want to list a file whose name contains 2 characters
ls ??

```
[hao@localhost sbin]$ ls ??  
ip  ss  tc
```

Range of character matching ([])

1. If we want to list files whose names start with 'a' or 'b'

`ls [ab]*`

```
[hao@localhost sbin]$ ls [ab]*
accessdb          audispd           badblocks         btrfs-convert
addgnupghome      auditctl          biosdecode        btrfs-debug-tree
addpart           auditd            biosdevname       btrfs-find-root
adduser           augenrules        blkdeactivate     btrfs-image
agetty            aureport          blkdiscard        btrfs-map-logical
alternatives      ausearch          blkid              btrfs-select-super
anacron           authconfig        blockdev          btrfstune
applygnupgdefaults authconfig-tui    bridge            btrfs-zero-log
arpd              autrace           btrfs              build-locale-
archive           avcstat           btrfsck
```


Range of character matching ([])

1. If we want to list files whose names contain numbers

`ls *[0-9]*`

```
[hao@localhost sbin]$ ls *[0-9]*
dumpe2fs          glibc_post_upgrade.x86_64  grub2-setpassword          mkfs.ext3
e2freefrag        grub2-bios-setup           grub2-sparc64-setup        mkfs.ext4
e2fsck            grub2-get-kernel-settings  iconvconfig.x86_64        pam_tally2
e2image           grub2-install              intel-microcode2ucode     ping6
e2label           grub2-mkconfig              ip6tables                  resize2fs
e2undo            grub2-ofpathname            ip6tables-restore
sasldblistusers2
e4defrag          grub2-probe                 ip6tables-save
fsck.ext2          grub2-reboot                killall5
fsck.ext3          grub2-rpm-sort              mke2fs
fsck.ext4          grub2-set-default           mkfs.ext2

```

Range of character matching ([])

1. If we want to list files whose names end with 'a-d'

`ls *[a-d]`

```
[hao@localhost sbin]$ ls *[a-d]
accessdb      dmfilemapd      lid             nl-class-add    thin_delta
arpd          firewallld      lpasswd         nl-cls-add      tuned
audispd       fxload          lsmod           nl-qdisc-add    unix_chkpwd
auditd        genhostid       luseradd        parted          useradd
blkdiscard    groupadd        lusermod        plymouthd       usermod
blkid         groupmod        lvextend        rdisc           vgextend
chpasswd      grub2-setpassword lvmetad         rmmod           xfs_db
chronyd       insmod          lvmpolld        rsyslogd        xfs_quota
crond         kexec          lvmsadc         selinuxenabled  zic
depmod        lgrouppadd      mkdumprd        sshd
dmeventd      lgrouppmod      mklost+found    tc
```

Regular Expression

- Regular expressions are standardized patterns that parse and manipulate text. For example, the regular expression:
 - I sent you a che(quelck) for the gr[ae]y-colou?red alumni?um
 - matches sentences that use either American or British spelling conventions.
- Regular expressions are supported by most modern languages.
- Wildcard is not a form of regular expression!

Special characters in regular expression

Symbol	What it matches or does
.	Matches any character
[]	Matches any character from a given set
^	Matches the beginning of a line
\$	Matches the end of a line
\w	Matches any "word" character
\s	Matches any whitespace character
\d	Matches any digits
	Matches either the element to its left or right
?	Allows zero or one match of the preceding element
*	Allows zero, one, or many matches of the preceding elements
+	Allows one or more matches of the preceding element
{n}	Matches exactly n instances of the preceding elements
{min,}	Matches at least min instance of
{min, max}	Matches any number of instances from min to max

Text Processing Tools: grep

- grep: print lines matching a pattern
- Syntax:
 - `grep [option] pattern [file]`
- Commonly used options:
 - i: ignore case distinctions in both the pattern and the input files.
 - n: Prefix each line of output with the line number within its input file
 - v: Invert the sense of matching, to select non-matching lines
- In basic regular expressions the meta-characters `?`, `+`, `{`, `|`, `(`, and `)` lose their special meaning; instead use the backslashed versions `\?`, `\+`, `\{`, `\|`, `\(`, and `\)`.

grep examples:

- We will first create a file for demonstration:
- `ps aux > grep.demo`
- If we want to find the lines contain 'root'
- `grep root grep.demo`

```
[hao@localhost ~]$ grep root grep.demo
root          1  0.0  0.6 128164  6824 ?        Ss   Feb26   0:01 /usr/lib/systemd/
systemd --switched-root --system --deserialize 21
root          2  0.0  0.0      0      0 ?        S    Feb26   0:00 [kthreadd]
root          3  0.0  0.0      0      0 ?        S    Feb26   0:00 [ksoftirqd/0]
root          5  0.0  0.0      0      0 ?        S<   Feb26   0:00 [kworker/0:0H]
root          6  0.0  0.0      0      0 ?        S    Feb26   0:00 [kworker/u2:0]
root          7  0.0  0.0      0      0 ?        S    Feb26   0:00 [migration/0]
root          8  0.0  0.0      0      0 ?        S    Feb26   0:00 [rcu_bh]
```

grep examples:

- If we want to find the lines do not contain 'root'
- `grep -v root grep.demo`

```
[hao@localhost ~]$ grep -v root grep.demo
USER          PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
polkitd        626  0.0  1.2 534892 12884 ?        Ssl  Feb26   0:00 /usr/lib/polkit-1/
polkitd --no-debug
dbus           627  0.0  0.1  32776  1856 ?        Ssl  Feb26   0:00 /bin/dbus-daemon --
system --address=systemd: --nofork --nopidfile --systemd-activation
chrony         635  0.0  0.1 115640  1780 ?        S    Feb26   0:00 /usr/sbin/chronyd
postfix        1113  0.0  0.3  89716  4028 ?        S    Feb26   0:00 qmgr -l -t unix -u
hao            2160  0.0  0.2 145700  2400 ?        S    Feb26   0:00 sshd: hao@pts/0
hao            2161  0.0  0.2 115392  2132 pts/0    Ss   Feb26   0:00 -bash
postfix        3031  0.0  0.3  89648  4004 ?        S    00:05   0:00 pickup -l -t unix -
u
hao            3070  0.0  0.1 151064  1808 pts/0    R+   00:15   0:00 ps aux
```

grep examples:

- If we want to find the lines contain 'root' and 'xfs'
- `grep root grep.demo | grep xfs`

```
[hao@localhost ~]$ grep root grep.demo | grep xfs
root      388  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfsalloc]
root      389  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs_mru_cache]
root      390  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-buf/dm-0]
root      391  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-data/dm-0]
root      392  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-conv/dm-0]
root      393  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-cil/dm-0]
root      394  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-reclaim/dm-]
root      395  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-log/dm-0]
root      396  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-eofblocks/d]
root      397  0.0  0.0      0      0 ?      S     Feb26  0:01  [xfsaield/dm-0]
root      543  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-buf/sda1]
root      546  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-data/sda1]
root      547  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-conv/sda1]
root      549  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-cil/sda1]
root      551  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-reclaim/sda]
root      553  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-log/sda1]
root      555  0.0  0.0      0      0 ?      S<    Feb26  0:00  [xfs-eofblocks/s]
root      566  0.0  0.0      0      0 ?      S     Feb26  0:00  [xfsaield/sda1]
```


grep examples:

- If we want to find the lines contain 'root' or 'xfs'
- `grep "root|xfs" grep.demo`

```
[hao@localhost ~]$ grep "root|xfs" grep.demo
root          1  0.0  0.6 128164  6824 ?          Ss    Feb26   0:01 /usr/lib/systemd/
systemd --switched-root --system --deserialize 21
root          2  0.0  0.0      0      0 ?          S     Feb26   0:00 [kthreadd]
root          3  0.0  0.0      0      0 ?          S     Feb26   0:00 [ksoftirqd/0]
root          5  0.0  0.0      0      0 ?        S<    Feb26   0:00 [kworker/0:0H]
root        389  0.0  0.0      0      0 ?        S<    Feb26   0:00 [xfs_mru_cache]
root        390  0.0  0.0      0      0 ?        S<    Feb26   0:00 [xfs-buf/dm-0]
root        391  0.0  0.0      0      0 ?        S<    Feb26   0:00 [xfs-data/dm-0]
root        392  0.0  0.0      0      0 ?        S<    Feb26   0:00 [xfs-conv/dm-0]
root        393  0.0  0.0      0      0 ?        S<    Feb26   0:00 [xfs-cil/dm-0]
root        394  0.0  0.0      0      0 ?        S<    Feb26   0:00 [xfs-reclaim/dm-]
root        395  0.0  0.0      0      0 ?        S<    Feb26   0:00 [xfs-log/dm-0]
```

grep examples:

- If we want to find the lines contain 'ss'
- `grep "ss" grep.demo`

```
[hao@localhost ~]$ grep ss grep.demo
dbus          627  0.0  0.1  32776  1856 ?          Ssl  Feb26   0:00 /bin/dbus-daemon --
system --adress=systemd: --nofork --nopidfile --systemd-activation
root          1009  0.0  0.4 105996  4072 ?          Ss   Feb26   0:00 /usr/sbin/sshd -D
root          2156  0.0  0.5 145700  5276 ?          Ss   Feb26   0:00 sshd: hao [priv]
hao           2160  0.0  0.2 145700  2400 ?          S    Feb26   0:00 sshd: hao@pts/0
```

grep examples:

- If we want to find the lines contain 'ss' or 'SS' or 'Ss' or 'sS'
- `grep -i "ss" grep.demo`

```
[hao@localhost ~]$ grep -i ss grep.demo
USER          PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root           1   0.0   0.6 128164  6824 ?        Ss   Feb26   0:01 /usr/lib/systemd/
systemd --switched-root --system --deserialize 21
root          464   0.0   0.3  34988  3788 ?        Ss   Feb26   0:00 /usr/lib/systemd/
systemd-journald
root          483   0.0   0.4 129552  4136 ?        Ss   Feb26   0:00 /usr/sbin/lvmेटad -f
root          495   0.0   0.4  46772  4824 ?        Ss   Feb26   0:00 /usr/lib/systemd/
systemd-udevd
root          624   0.0   0.1  24204  1708 ?        Ss   Feb26   0:00 /usr/lib/systemd/
systemd-logind
polkitd       626   0.0   1.2 534892 12884 ?        Ssl  Feb26   0:00 /usr/lib/polkit-1/
polkitd --no-debug
```

grep examples:

- If we want to find the lines contain 'ps' from multiple files:
- We make another file called grep.demo1 by using command:
 - `ps > grep.demo1`
 - `grep ps grep.demo grep.demo1`

```
[hao@localhost ~]$ grep ps grep.demo grep.demo1
grep.demo:root          39  0.0  0.0          0      0 ?          S<    Feb26
0:00 [kpsmoused]
grep.demo:hao          3070  0.0  0.1 151064   1808 pts/0      R+    00:15
0:00 ps aux
grep.demo1: 3180 pts/0    00:00:00 ps
```

Using grep and regular expression

- grep command is usually been used with regular expression.
- If we want to list all the processes started in Feb:
grep 'Feb..' grep.demo

```
[hao@localhost ~]$ grep 'Feb..' grep.demo
root          1  0.0  0.6 128164  6824 ?          Ss   Feb26    0:01 /usr/lib/
systemd/systemd --switched-root --system --deserialize 21
root          2  0.0  0.0      0      0 ?          S    Feb26    0:00 [kthreadd]
root          3  0.0  0.0      0      0 ?          S    Feb26    0:00 [ksoftirqd/0]
root          5  0.0  0.0      0      0 ?        S<   Feb26    0:00 [kworker/0:0H]
root          6  0.0  0.0      0      0 ?          S    Feb26    0:00 [kworker/u2:0]
root          7  0.0  0.0      0      0 ?          S    Feb26    0:00 [migration/0]
```

Using grep and regular expression

- grep command is usually been used with regular expression.
- If we want to list all the lines start with p :
- `grep '^p' grep.demo`

```
[hao@localhost ~]$ grep '^p' grep.demo
p polkitd      626  0.0  1.2 534892 12884 ?          Ssl  Feb26   0:00 /usr/lib/polkit-1/
p polkitd --no-debug
p postfix     1113  0.0  0.3  89716  4028 ?          S    Feb26   0:00 qmgr -l -t unix -u
p postfix     3031  0.0  0.3  89648  4004 ?          S    00:05   0:00 pickup -l -t unix -
u
```

Using grep and regular expression

- grep command is usually been used with regular expression.
- If we want to list all the lines end with a number :
- `grep '[0-9]$(grep.demo`

```
[hao@localhost ~]$ grep '^p' grep.demo
polkitd      626  0.0  1.2 534892 12884 ?          Ssl  Feb26   0:00 /usr/lib/polkit-1/
polkitd --no-debug
postfix      1113  0.0  0.3  89716  4028 ?          S    Feb26   0:00 qmgr -l -t unix -u
postfix      3031  0.0  0.3  89648  4004 ?          S    00:05   0:00 pickup -l -t unix -u
[hao@localhost ~]$ grep '[0-9]$(grep.demo
root          1  0.0  0.6 128164  6824 ?          Ss   Feb26   0:01 /usr/lib/systemd/
systemd --switched-root --system --deserialize 21
root         1921  0.0  1.5 113372 15912 ?          S    Feb26   0:00 /sbin/dhclient -d -q -
sf /usr/libexec/nm-dhcp-helper -pf /var/run/dhclient-enp0s3.pid -lf /var/lib/
NetworkManager/dhclient-dd33654a-838f-4664-abb4-38cf1b5980ac-enp0s3.lease -cf /var/lib/
NetworkManager/dhclient-enp0s3.conf enp0s3
root         2101  0.0  1.5 113372 15896 ?          S    Feb26   0:00 /sbin/dhclient -d -q -
sf /usr/libexec/nm-dhcp-helper -pf /var/run/dhclient-enp0s8.pid -lf /var/lib/
NetworkManager/dhclient-fad5ea84-d774-3cc6-a5b0-842365440f08-enp0s8.lease -cf /var/lib/
NetworkManager/dhclient-enp0s8.conf enp0s8
hao          2160  0.0  0.2 145700  2400 ?          S    Feb26   0:00 sshd: hao@pts/0
```

Exercise

- Please create a file contain the following information:

ID	Name	City	Advisor	DOB	Year	GPA
282102	Andy	New Haven	Lisa	01/23/99	Sr	3.5
281734	Zoe	Bridgeport	Hao	05/14/00	Jr	3.2
292342	Alex	Cheshire	hao	12/25/01	Fresh	2.0
284323	Eric	Hartford	lisa	02/26/00	Soph	2.5
274342	Bill	Orange	Imad	04/12/99	Sr	4.0
293232	Alice	New Haven	IMAD	07/04/98	Jr	1.8
213852	Lee	bridgeport	lisa	06/22/02	Fresh	3.8

Questions:

- Find the students whose advisor is Lisa
- Find the students come from orange
- Find the students come from New Haven
- Find the students do not come from New Haven
- Find the students born after 2000
- Find the students have GPA higher than 3
- Find the students whose advisor is Hao and have GPA higher than 3
- Find the students with their ID start with 28

awk: text processing and data extraction tool

- awk: is a language for processing text files
- Syntax:
 - `awk [options] 'command' input_file`
- Commonly used options:
 - F: field separator
- commands:
 - build-in functions are enclosed in {}

awk: text processing and data extraction tool

- Example: separate field in /etc/passwd file and print out username:
- `awk -F: '{print "username:" $1}' /etc/passwd`
- `$1`: the first field in the line
- `$0`: the whole line

```
[root@localhost ~]# awk -F: '{print "username:" $1}' /etc/passwd
username:root
username:bin
username:daemon
username:adm
username:lp
username:sync
username:shutdown
```

awk with regular expression

- Example: separate field in /etc/passwd file and print out root :
- `awk -F: '/root/{print "username:" $1}' /etc/passwd`

```
[root@localhost ~]# awk -F: '/root/{print "username:" $1}' /etc/passwd
username:root
username:operator
```

awk with regular expression

- Example: separate field in /etc/passwd file and print out root :
- we only want first line
- `awk -F: '$1=="root"{print "username:" $1}' /etc/passwd`

```
[root@localhost ~]# awk -F: '$1=="root"{print "username:" $1}' /etc/passwd
username:root
```

Exercise:

- Print the students name whose advisor is Lisa
- Print the name and DOB of students who come from orange
- Print the name and GPA of students who come from New Haven
- Print the name and hometown of students do not come from New Haven
- Print the DOB of students born after 2000
- Print the name, GPA, and advisor of students have GPA higher than 3
- Print the GPA and name of students whose advisor is Hao and have GPA higher than 3
- Print the name and year of students with GPA less than 3

Solution for Exercise 1:

```
grep -i lisa students
grep -i orange students
grep "New Haven" students
grep -v "New Haven" students
grep "[0-9][0-9]/[0-9][0-9]/0[0-9]" students
grep "[3-9]\.[0-9]" students
grep "[3-9]\.[0-9]" students|grep -i hao
grep "^28" students
```

Solution for Exercise 2:

```
awk '/[Ll]isa/{print $2}' students
```

I can also print out the header before names:

```
head -1 students|awk '{print $2}';awk '/[lL]isa/{print $2}' students
```

```
awk '/Orange/{print $2" " $5}' students
```

```
awk '/New Haven/{print $2" " $8}' students
```

```
grep -v "New Haven" students|awk '{print $2" "$3}'
```

```
awk '/[0-9][0-9]\/[0-9][0-9]\/[0-9][0-9]/{print $5}' students
```

```
awk '/[34]\.[0-9]$/ {if ($4=="Haven") print $2" "$5" "$8; else print $2" "$4" "$7}' students
```

```
awk '/[34]\.[0-9]$/ {if ($4=="Hao" || $4=="hao") print $7" "$2;} ' students
```

```
awk '/[0-2]\.[0-9]$/ {if ($4=="Haven") print $2" "$6; else print $2" "$5}' students
```