

CSC424 System Administration

Instructor: Dr. Hao Wu

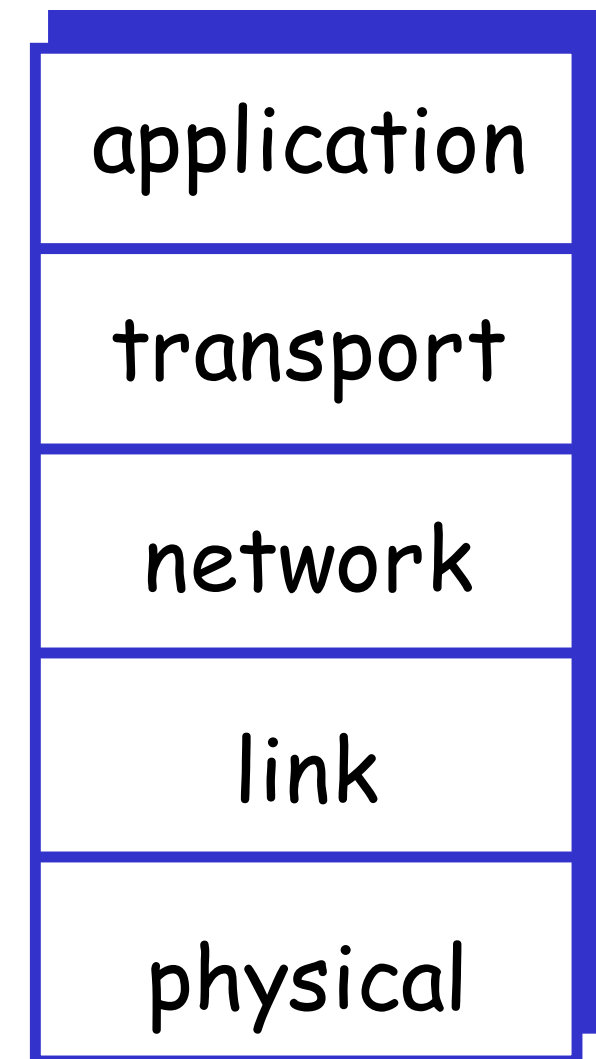
Week 11 Network and System Monitoring

Networking Basics

- Internet protocol layers
- MAC
- IP
- TCP/UDP
- Routing

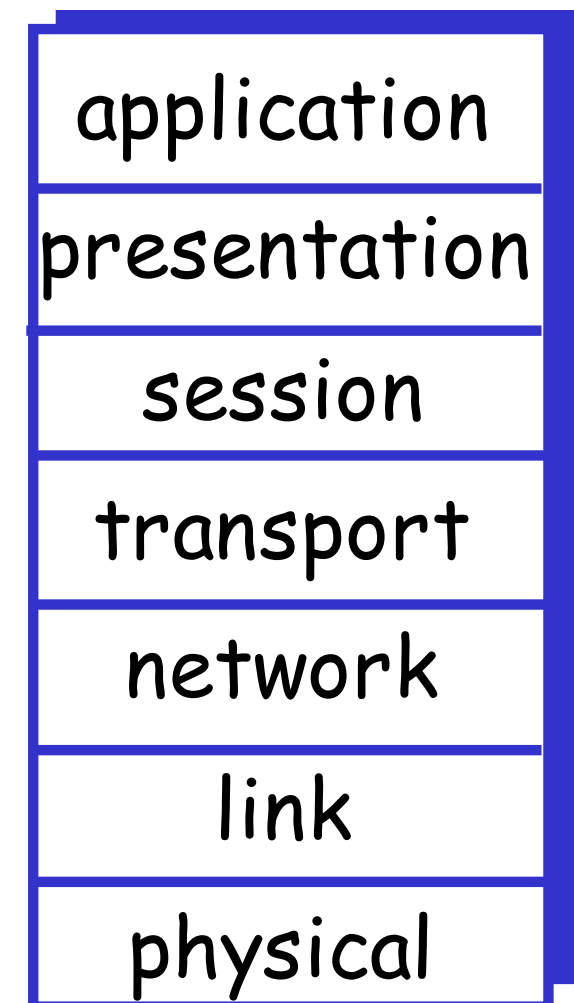
Internet protocol stack

- ❖ *application*: supporting network applications
 - FTP, SMTP, HTTP
- ❖ *transport*: process-process data transfer
 - TCP, UDP
- ❖ *network*: routing of datagrams from source to destination
 - IP, routing protocols
- ❖ *link*: data transfer between neighboring network elements
 - Ethernet, 802.111 (WiFi), PPP
- ❖ *physical*: bits “on the wire”

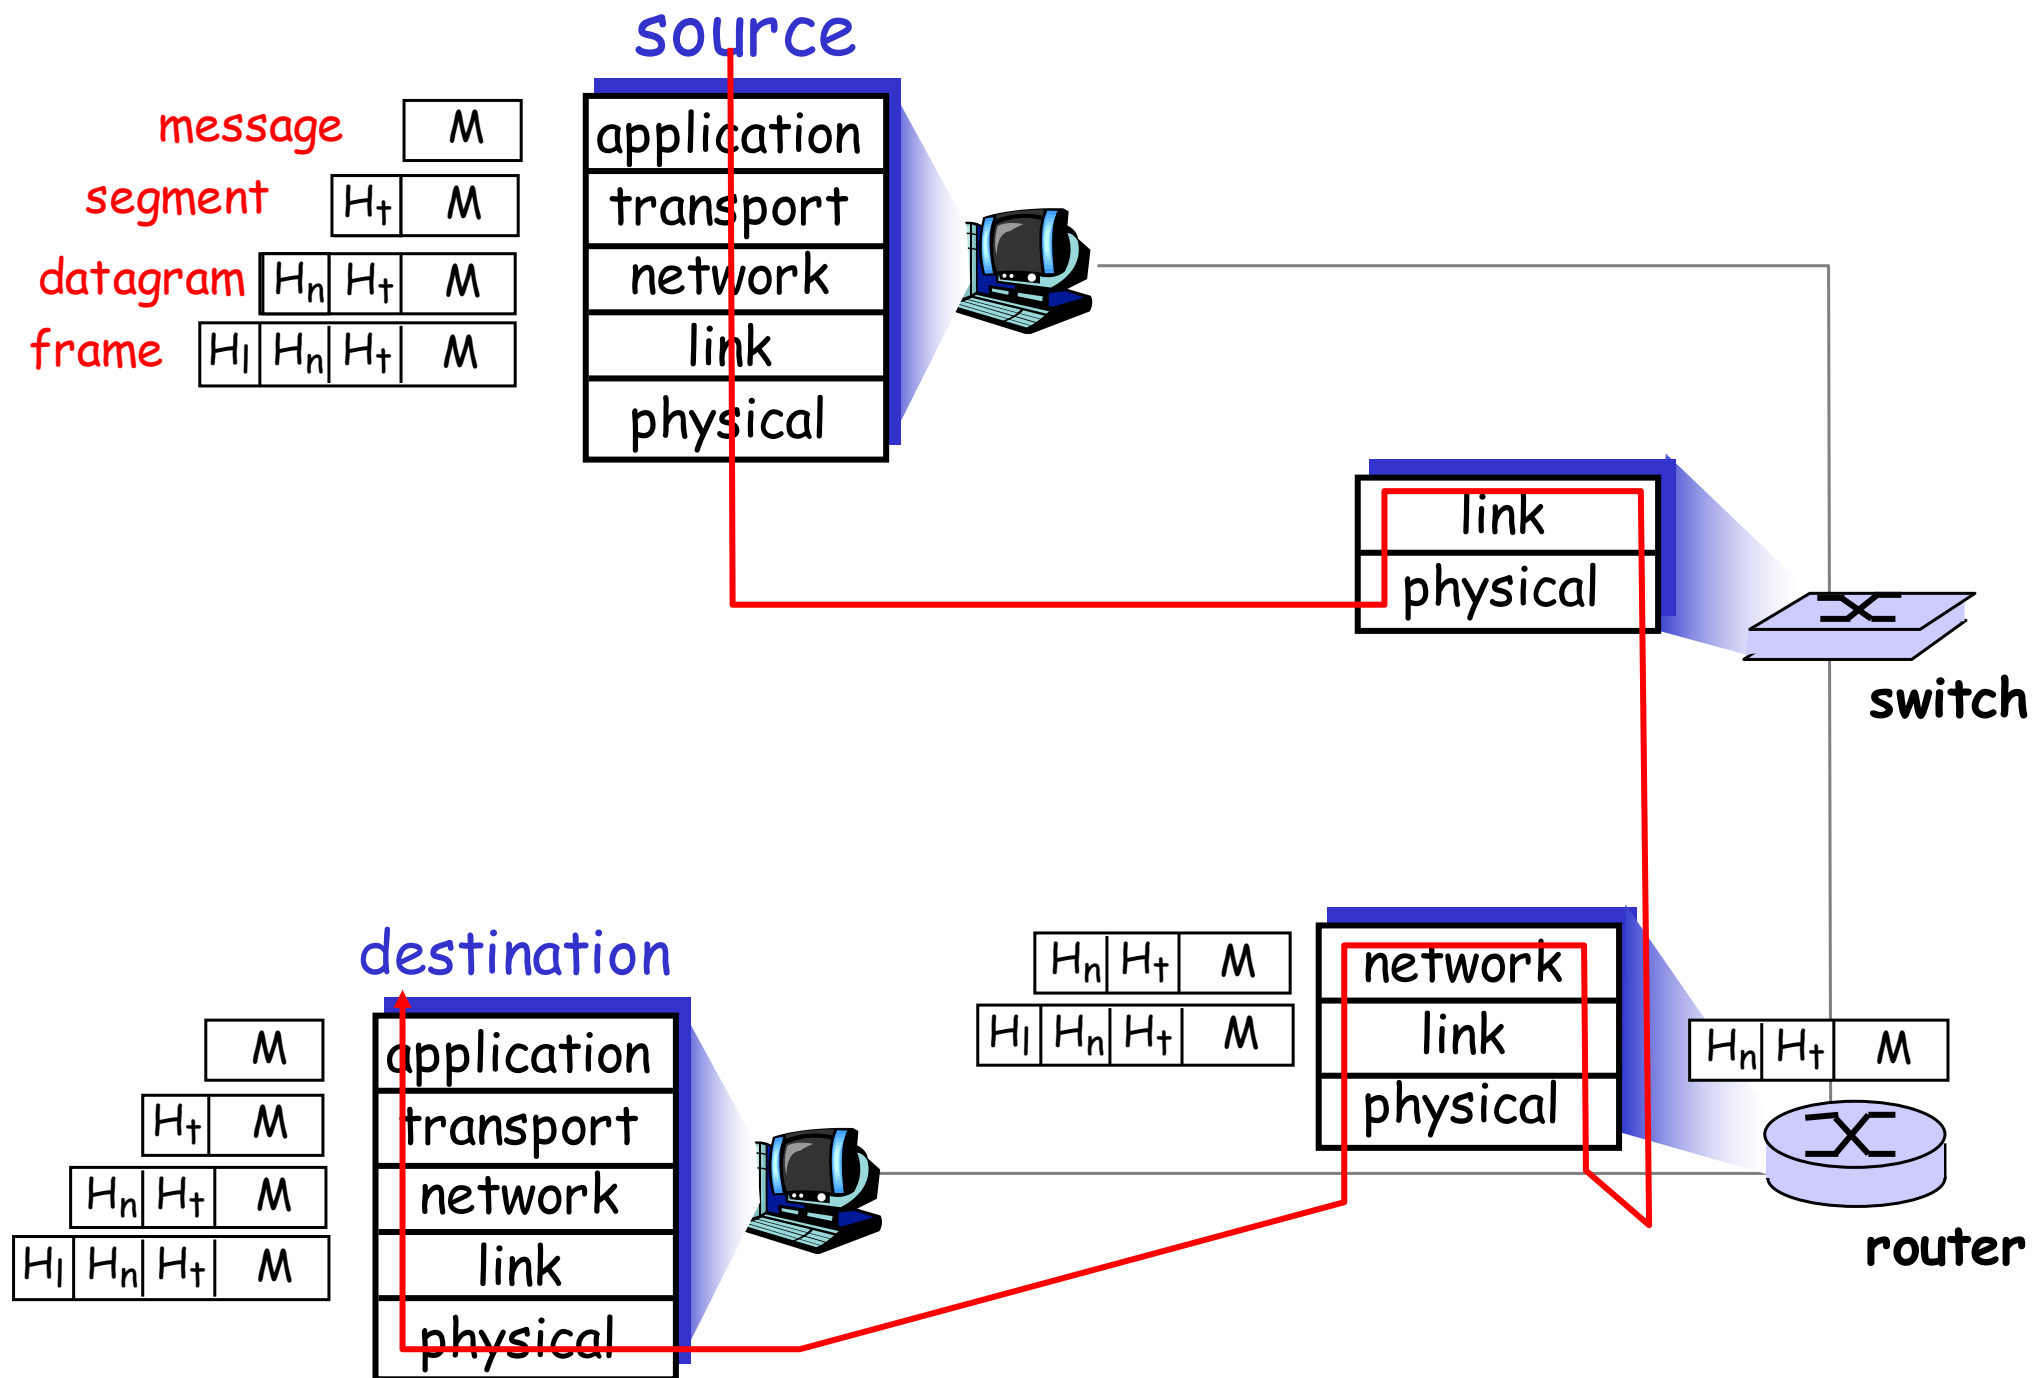


ISO/OSI (Open Systems Interconnection) reference model

- ❑ **presentation:** allow applications to interpret meaning of data, e.g., encryption, compression, machine-specific conventions
- ❑ **session:** synchronization, checkpointing, recovery of data exchange
- ❑ Internet stack “missing” these layers!
 - ❖ these services, *if needed*, must be implemented in application
 - ❖ needed?



Encapsulation

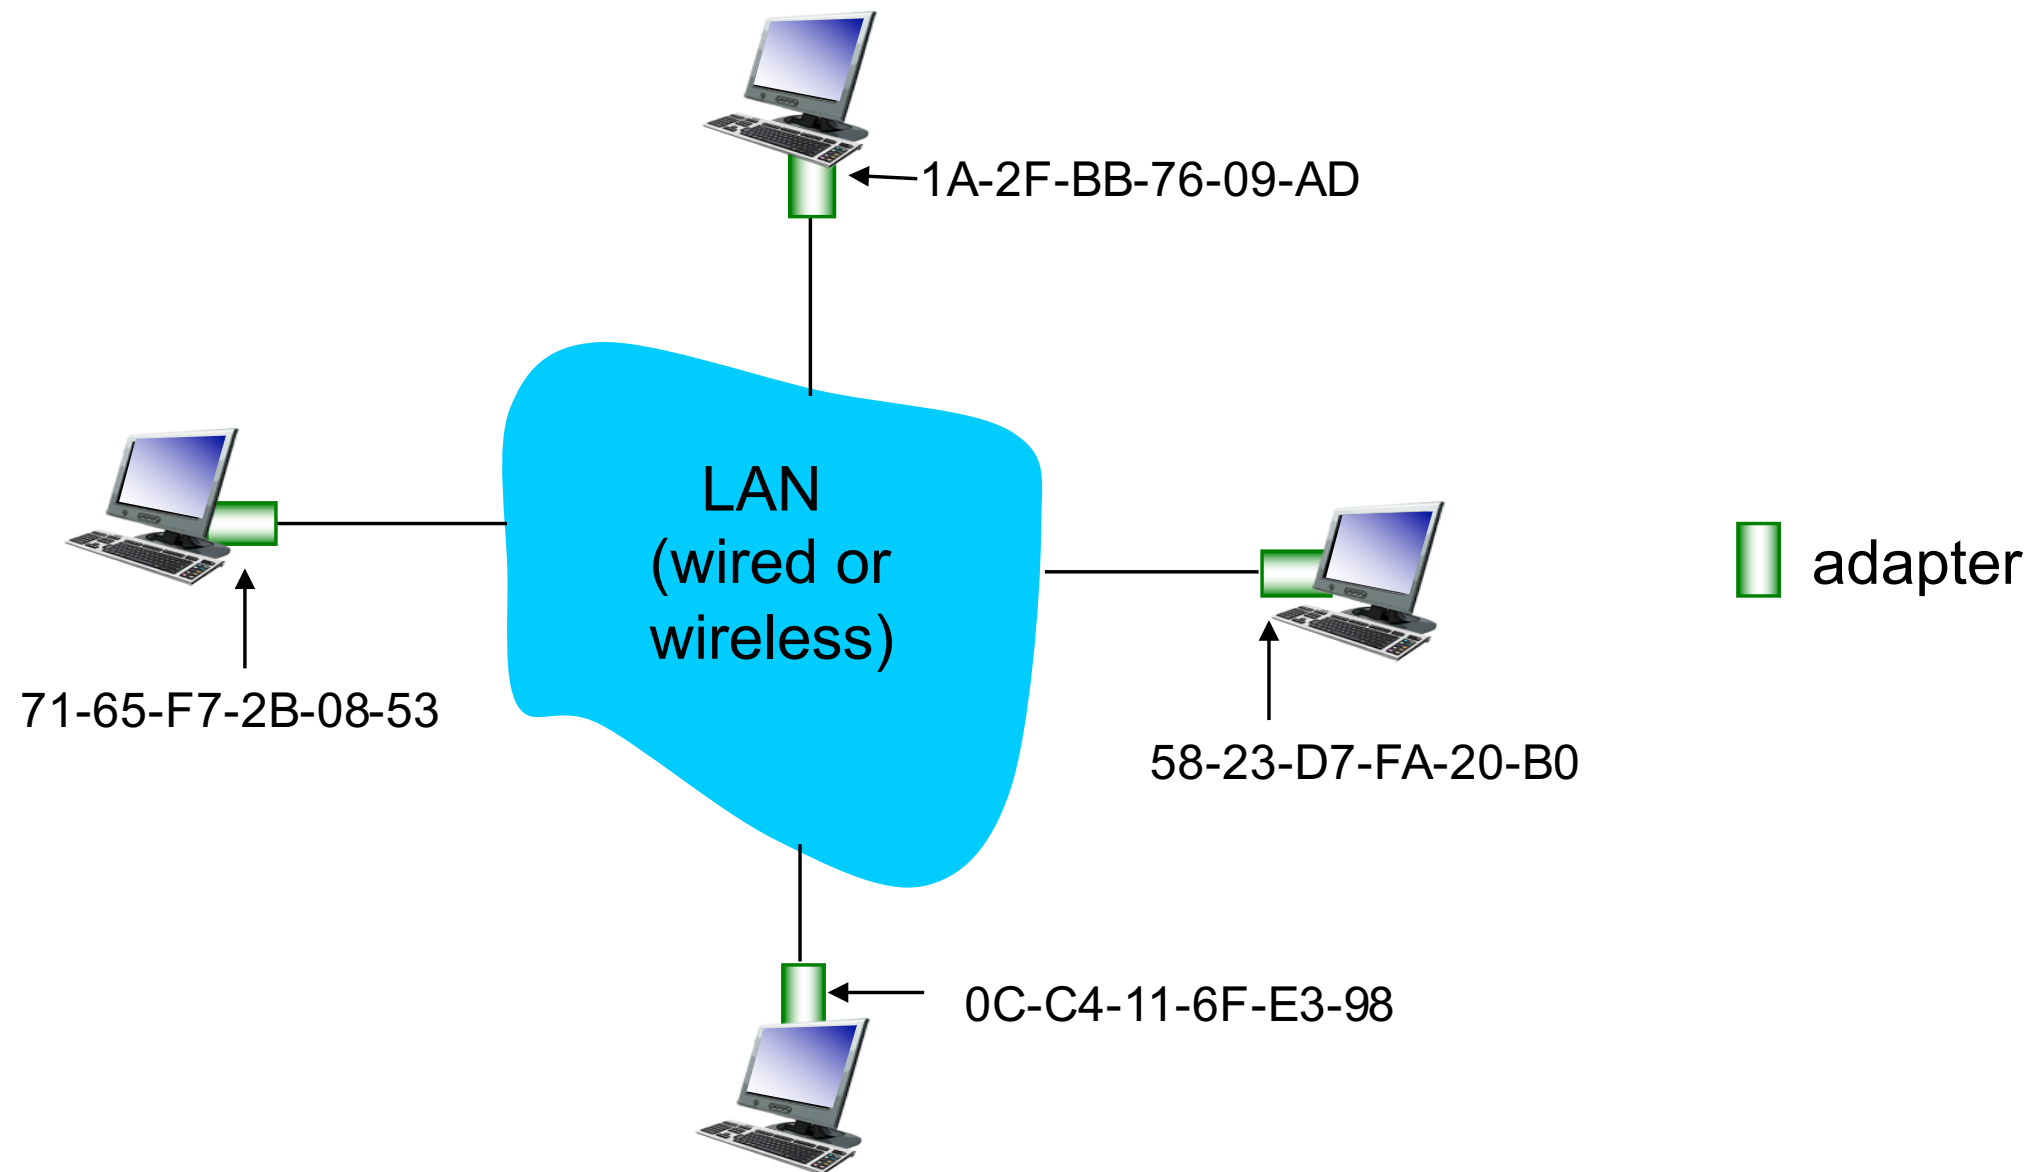


MAC addresses and ARP

- ❖ 32-bit IP address:
 - *network-layer* address for interface
 - used for layer 3 (network layer) forwarding
- ❖ MAC (or LAN or physical or Ethernet) address:
 - function: *used 'locally' to get frame from one interface to another physically-connected interface (same network, in IP-addressing sense)*
 - 48 bit MAC address (for most LANs) burned in NIC ROM, also sometimes software settable
 - e.g.: 1A-2F-BB-76-09-AD

LAN addresses and ARP

each adapter on LAN has unique **LAN** address



LAN addresses (more)

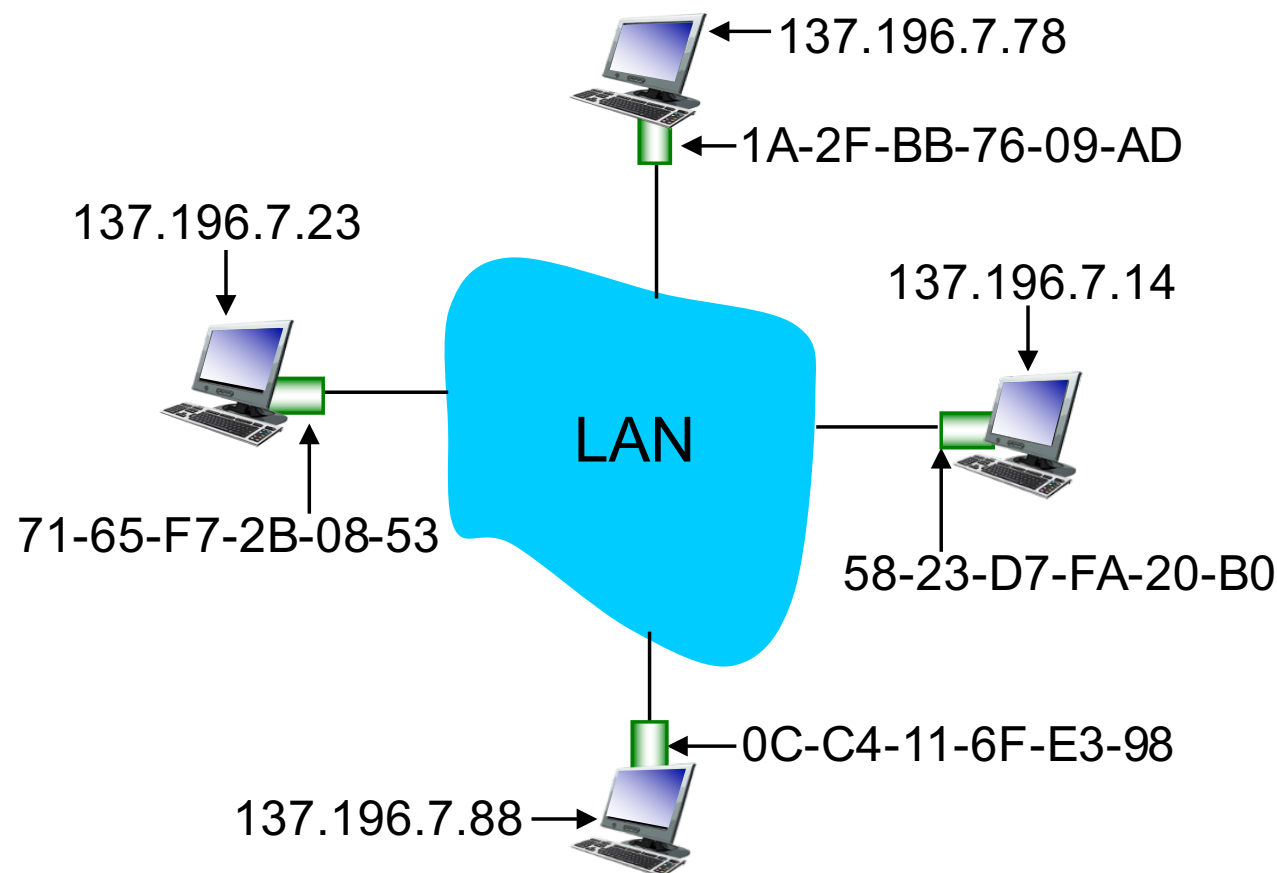
- ❖ MAC address allocation administered by IEEE
- ❖ manufacturer buys portion of MAC address space (to assure uniqueness)
- ❖ analogy:
 - MAC address: like Social Security Number
 - IP address: like postal address
- ❖ MAC flat address → portability
 - can move LAN card from one LAN to another
- ❖ IP hierarchical address *not* portable
 - address depends on IP subnet to which node is attached

ARP: address resolution protocol

Question: how to determine interface's MAC address, knowing its IP address?

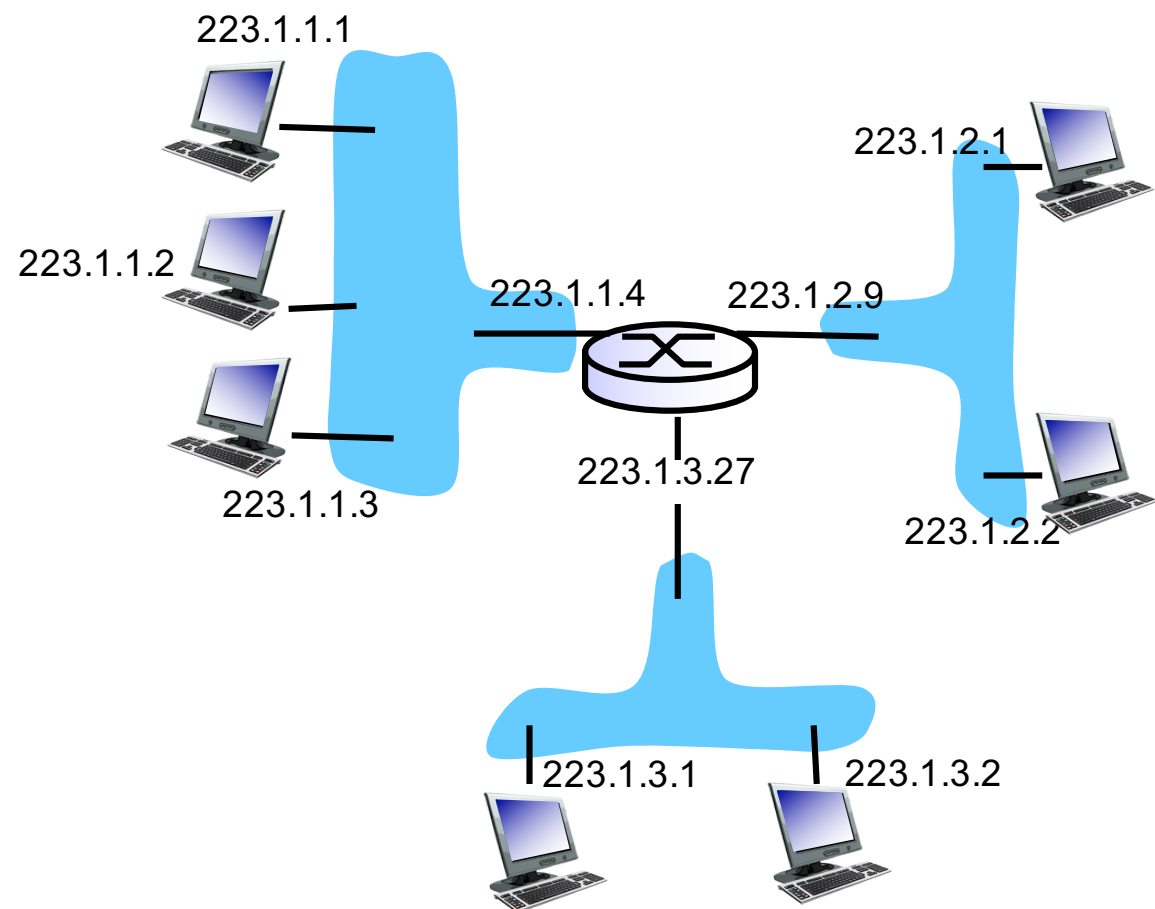
ARP table: each IP node (host, router) on LAN has table

- IP/MAC address mappings for some LAN nodes:
< IP address; MAC address; TTL >
- TTL (Time To Live): time after which address mapping will be forgotten (typically 20 min)



IP addressing: introduction

- ❖ **IP address:** 32-bit identifier for host, router *interface*
- ❖ **interface:** connection between host/router and physical link
 - router's typically have multiple interfaces
 - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)
- ❖ **IP addresses associated with each interface**



$$223.1.1.1 = \underbrace{11011111}_{223} \underbrace{00000001}_1 \underbrace{00000001}_1 \underbrace{00000001}_1$$

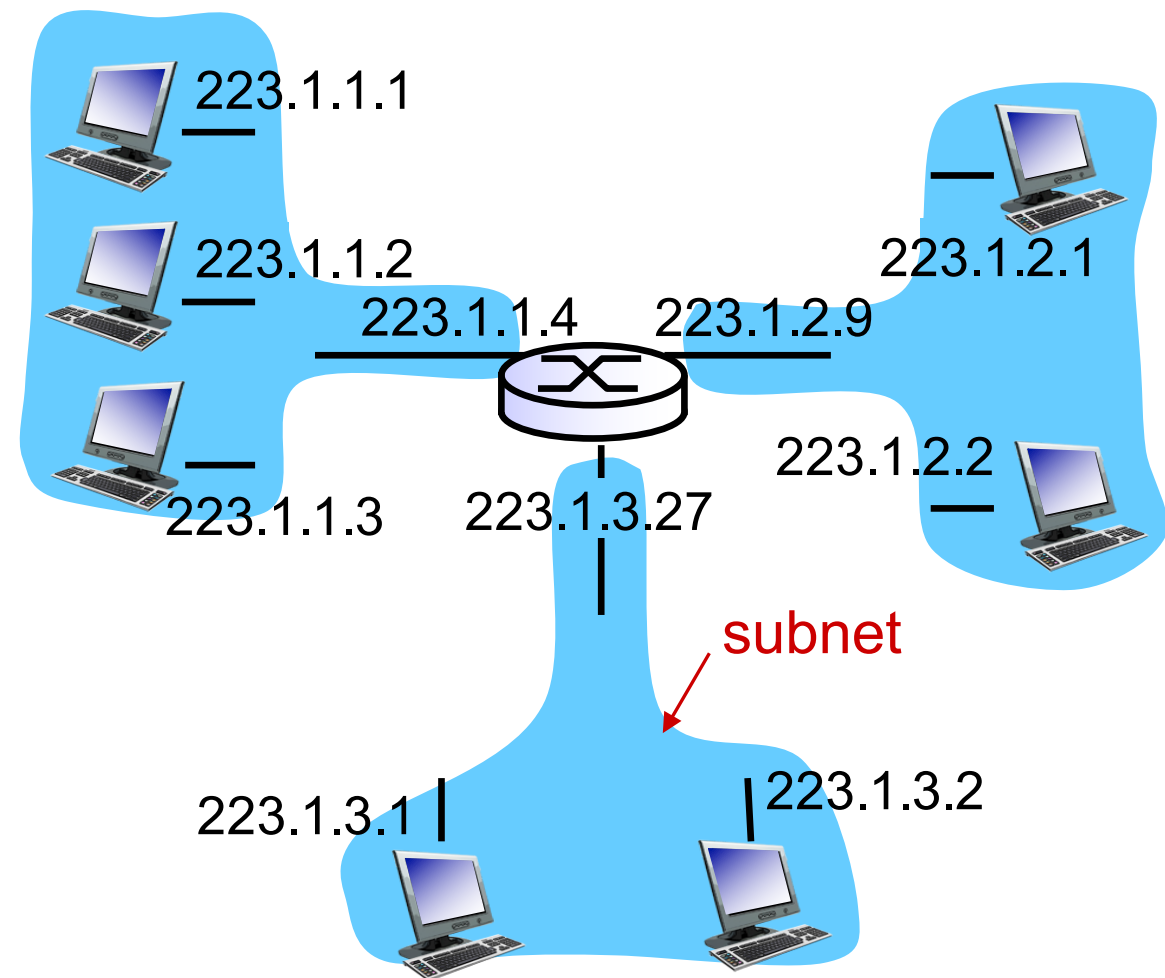
Subnets

❖ IP address:

- subnet part - high order bits
- host part - low order bits

❖ *what's a subnet?*

- device interfaces with same subnet part of IP address
- can physically reach each other *without intervening router*

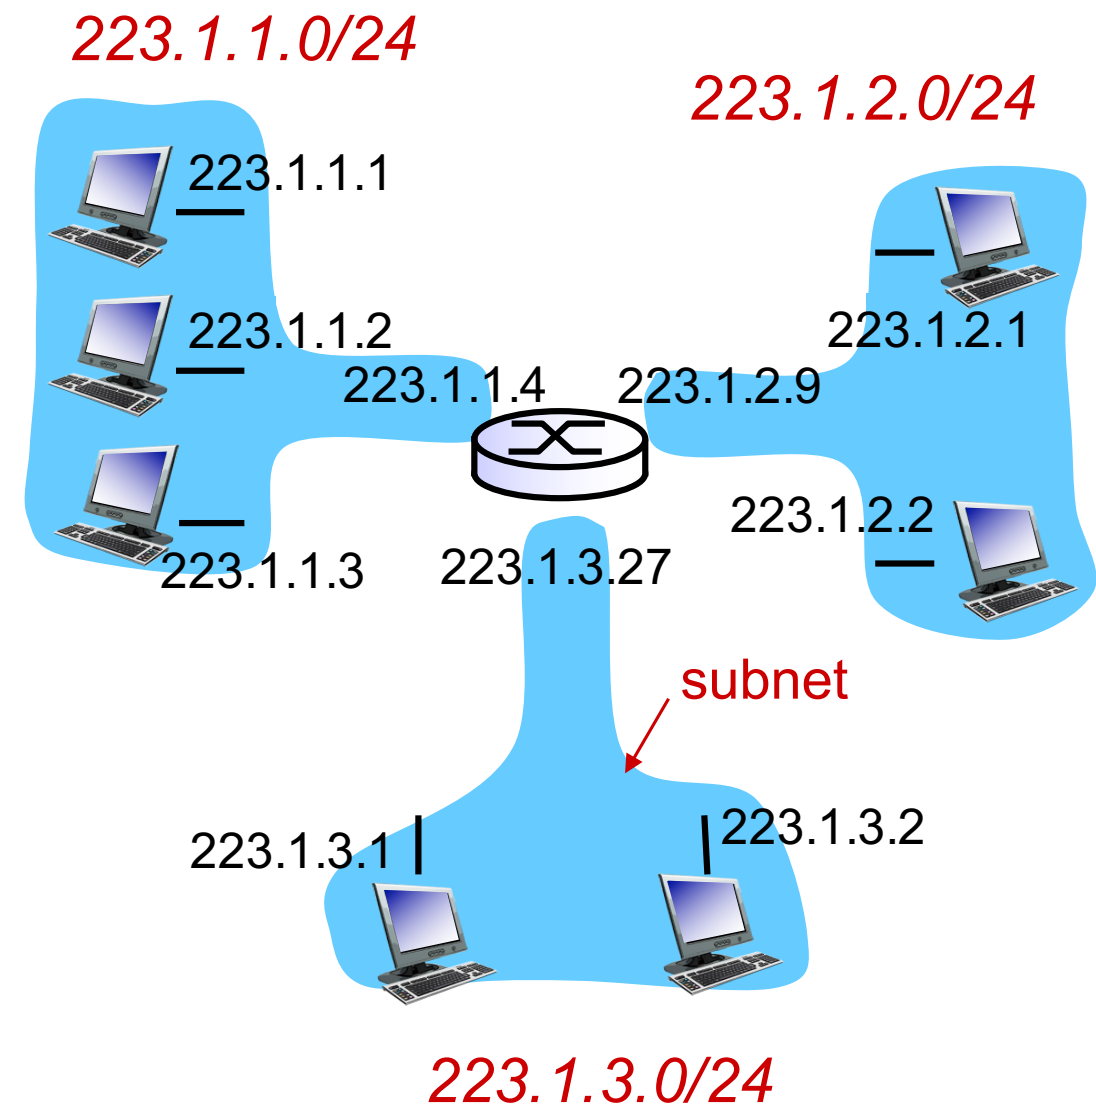


network consisting of 3 subnets

Subnets

recipe

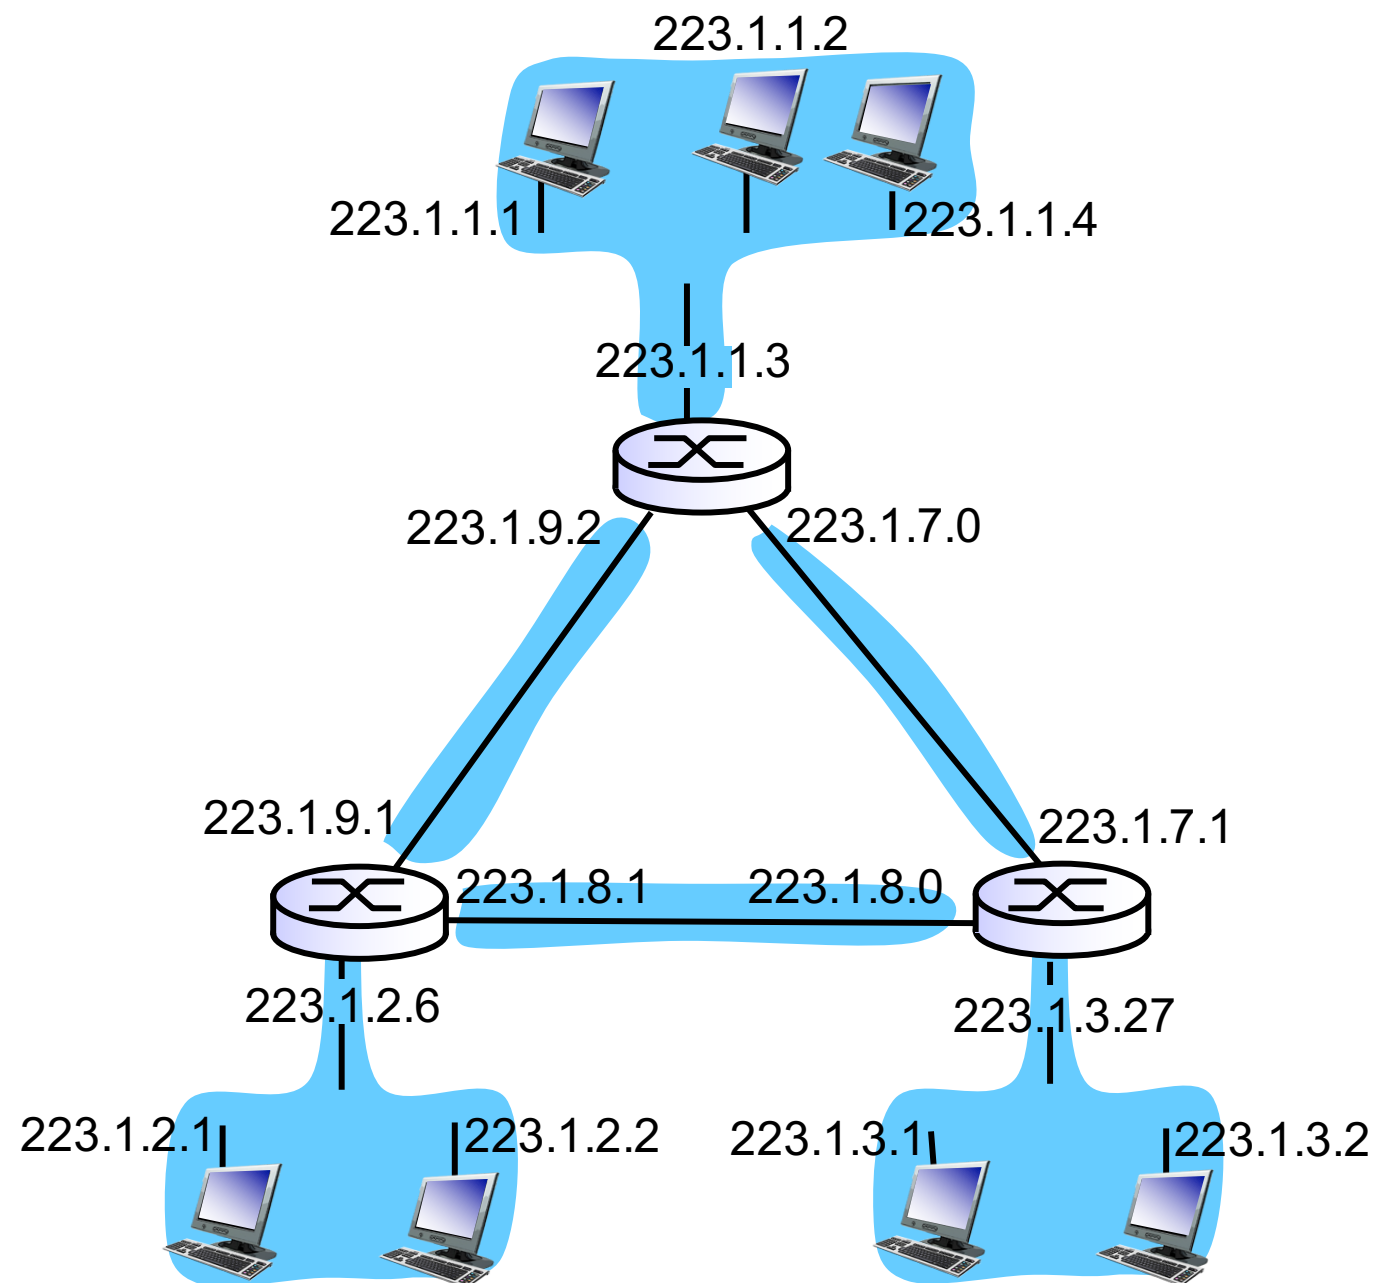
- ❖ to determine the subnets, detach each interface from its host or router, creating islands of isolated networks
- ❖ each isolated network is called a *subnet*



subnet mask: /24

Subnets

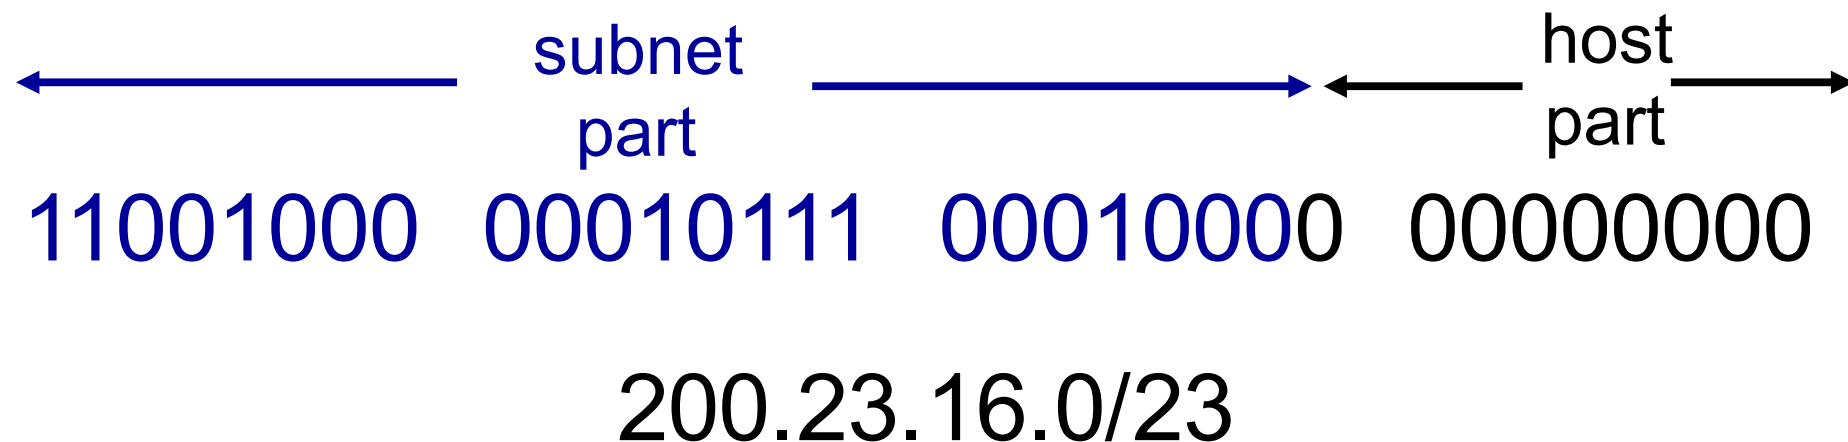
how many?



IP addressing: CIDR

CIDR: Classless InterDomain Routing

- subnet portion of address of arbitrary length
- address format: **a.b.c.d/x**, where x is # bits in subnet portion of address



IP addresses: how to get one?

Q: How does a *host* get IP address?

- ❖ hard-coded by system admin in a file
 - Windows: control-panel->network->configuration->tcp/ip->properties
 - UNIX: /etc/rc.config
- ❖ **DHCP: Dynamic Host Configuration Protocol:** dynamically get address from as server
 - “plug-and-play”

DHCP: Dynamic Host Configuration Protocol

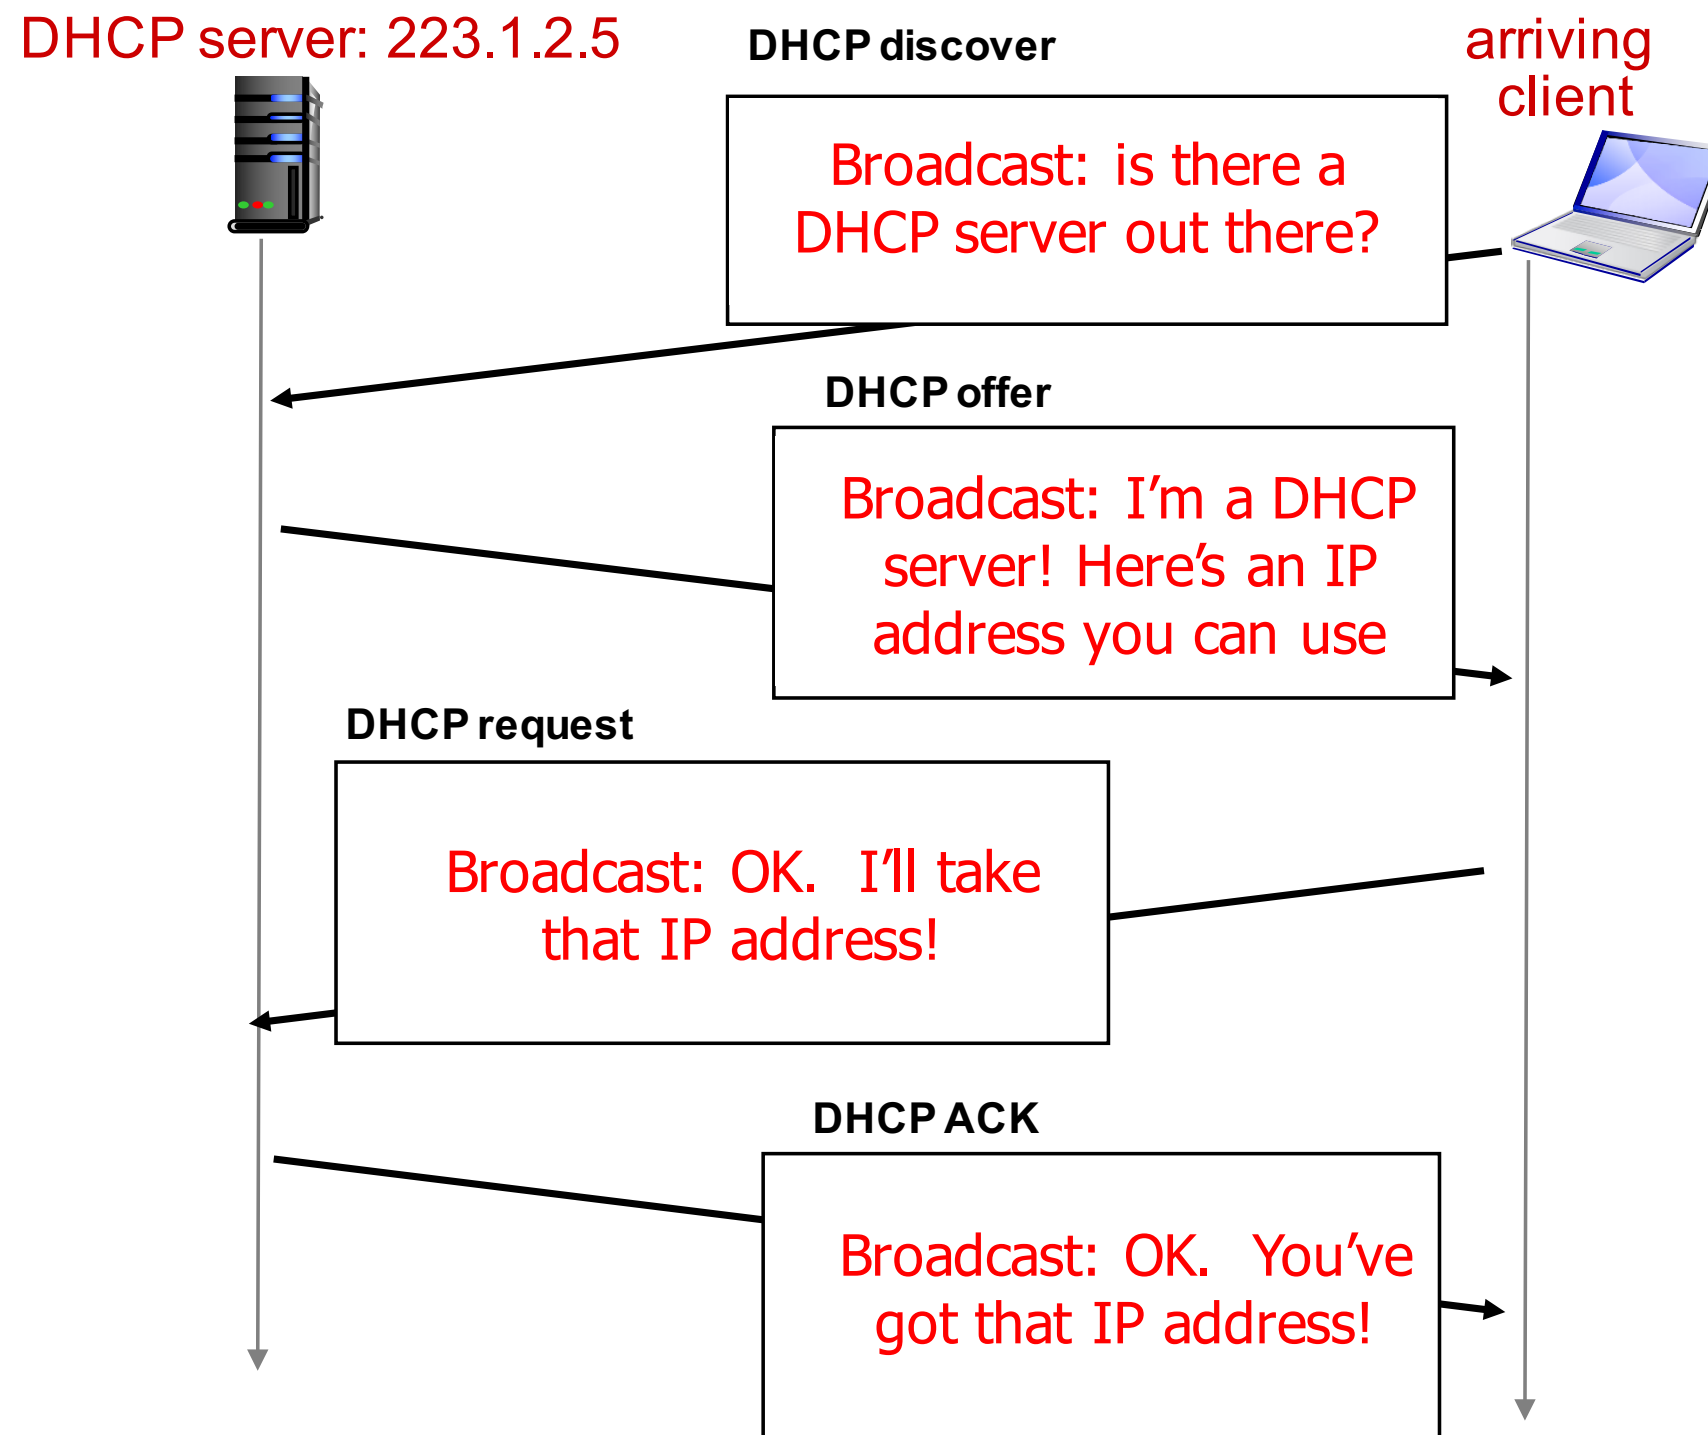
goal: allow host to *dynamically* obtain its IP address from network server when it joins network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/“on”)
- support for mobile users who want to join network (more shortly)

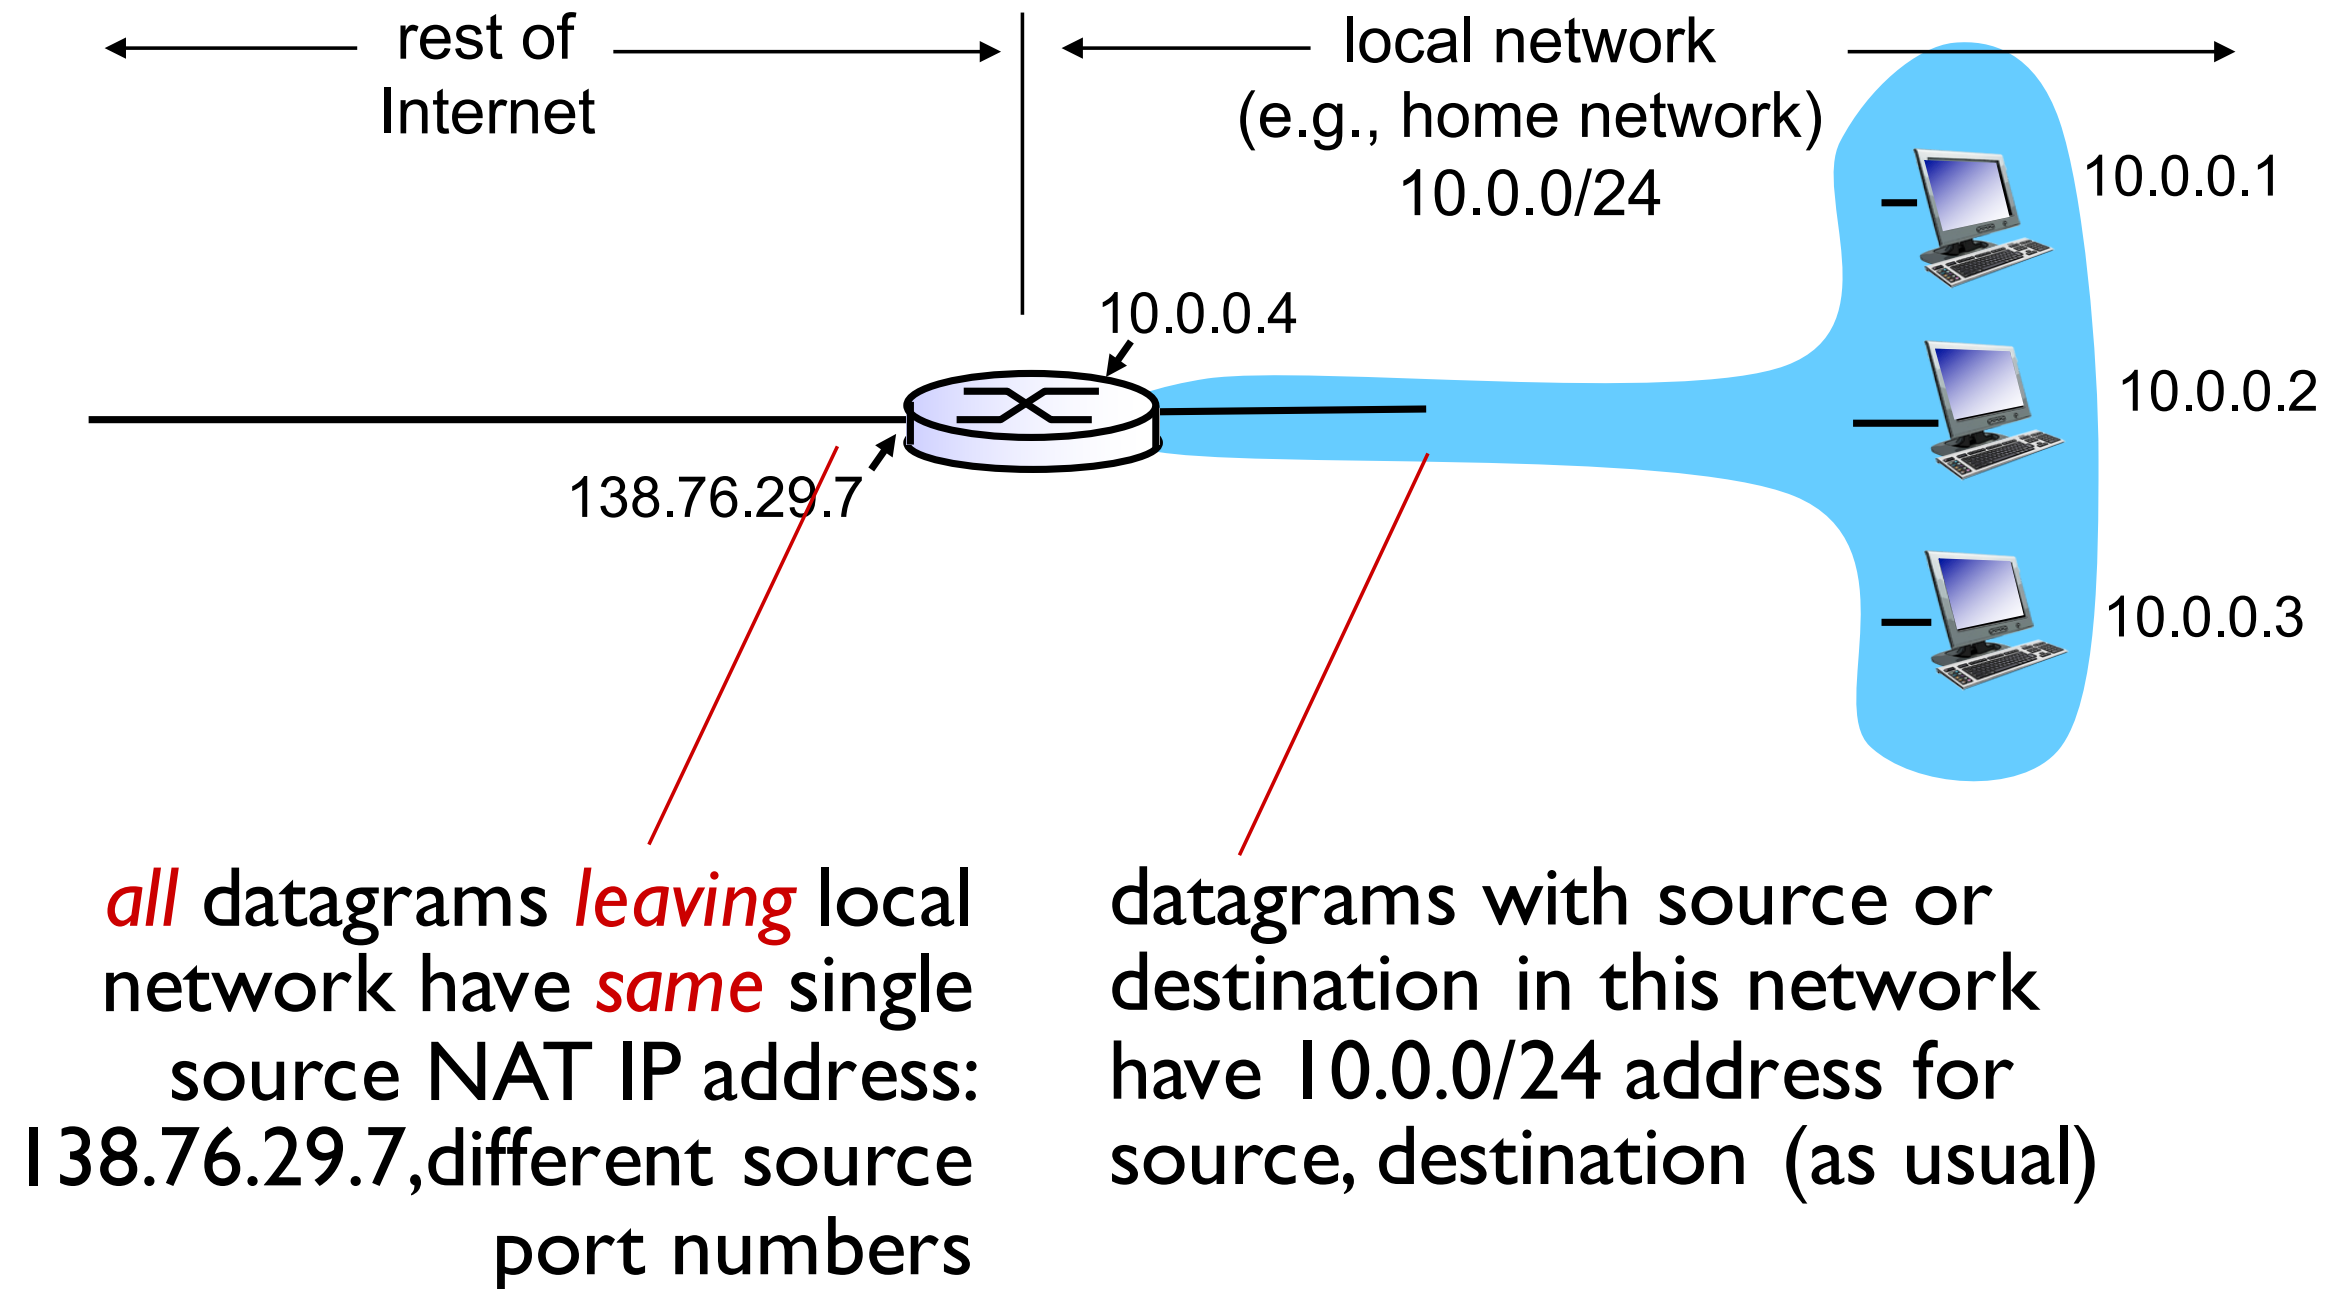
DHCP overview:

- host broadcasts “**DHCP discover**” msg [optional]
- DHCP server responds with “**DHCP offer**” msg [optional]
- host requests IP address: “**DHCP request**” msg
- DHCP server sends address: “**DHCP ack**” msg

DHCP client-server scenario



NAT: network address translation



NAT: network address translation

motivation: local network uses just one IP address as far as outside world is concerned:

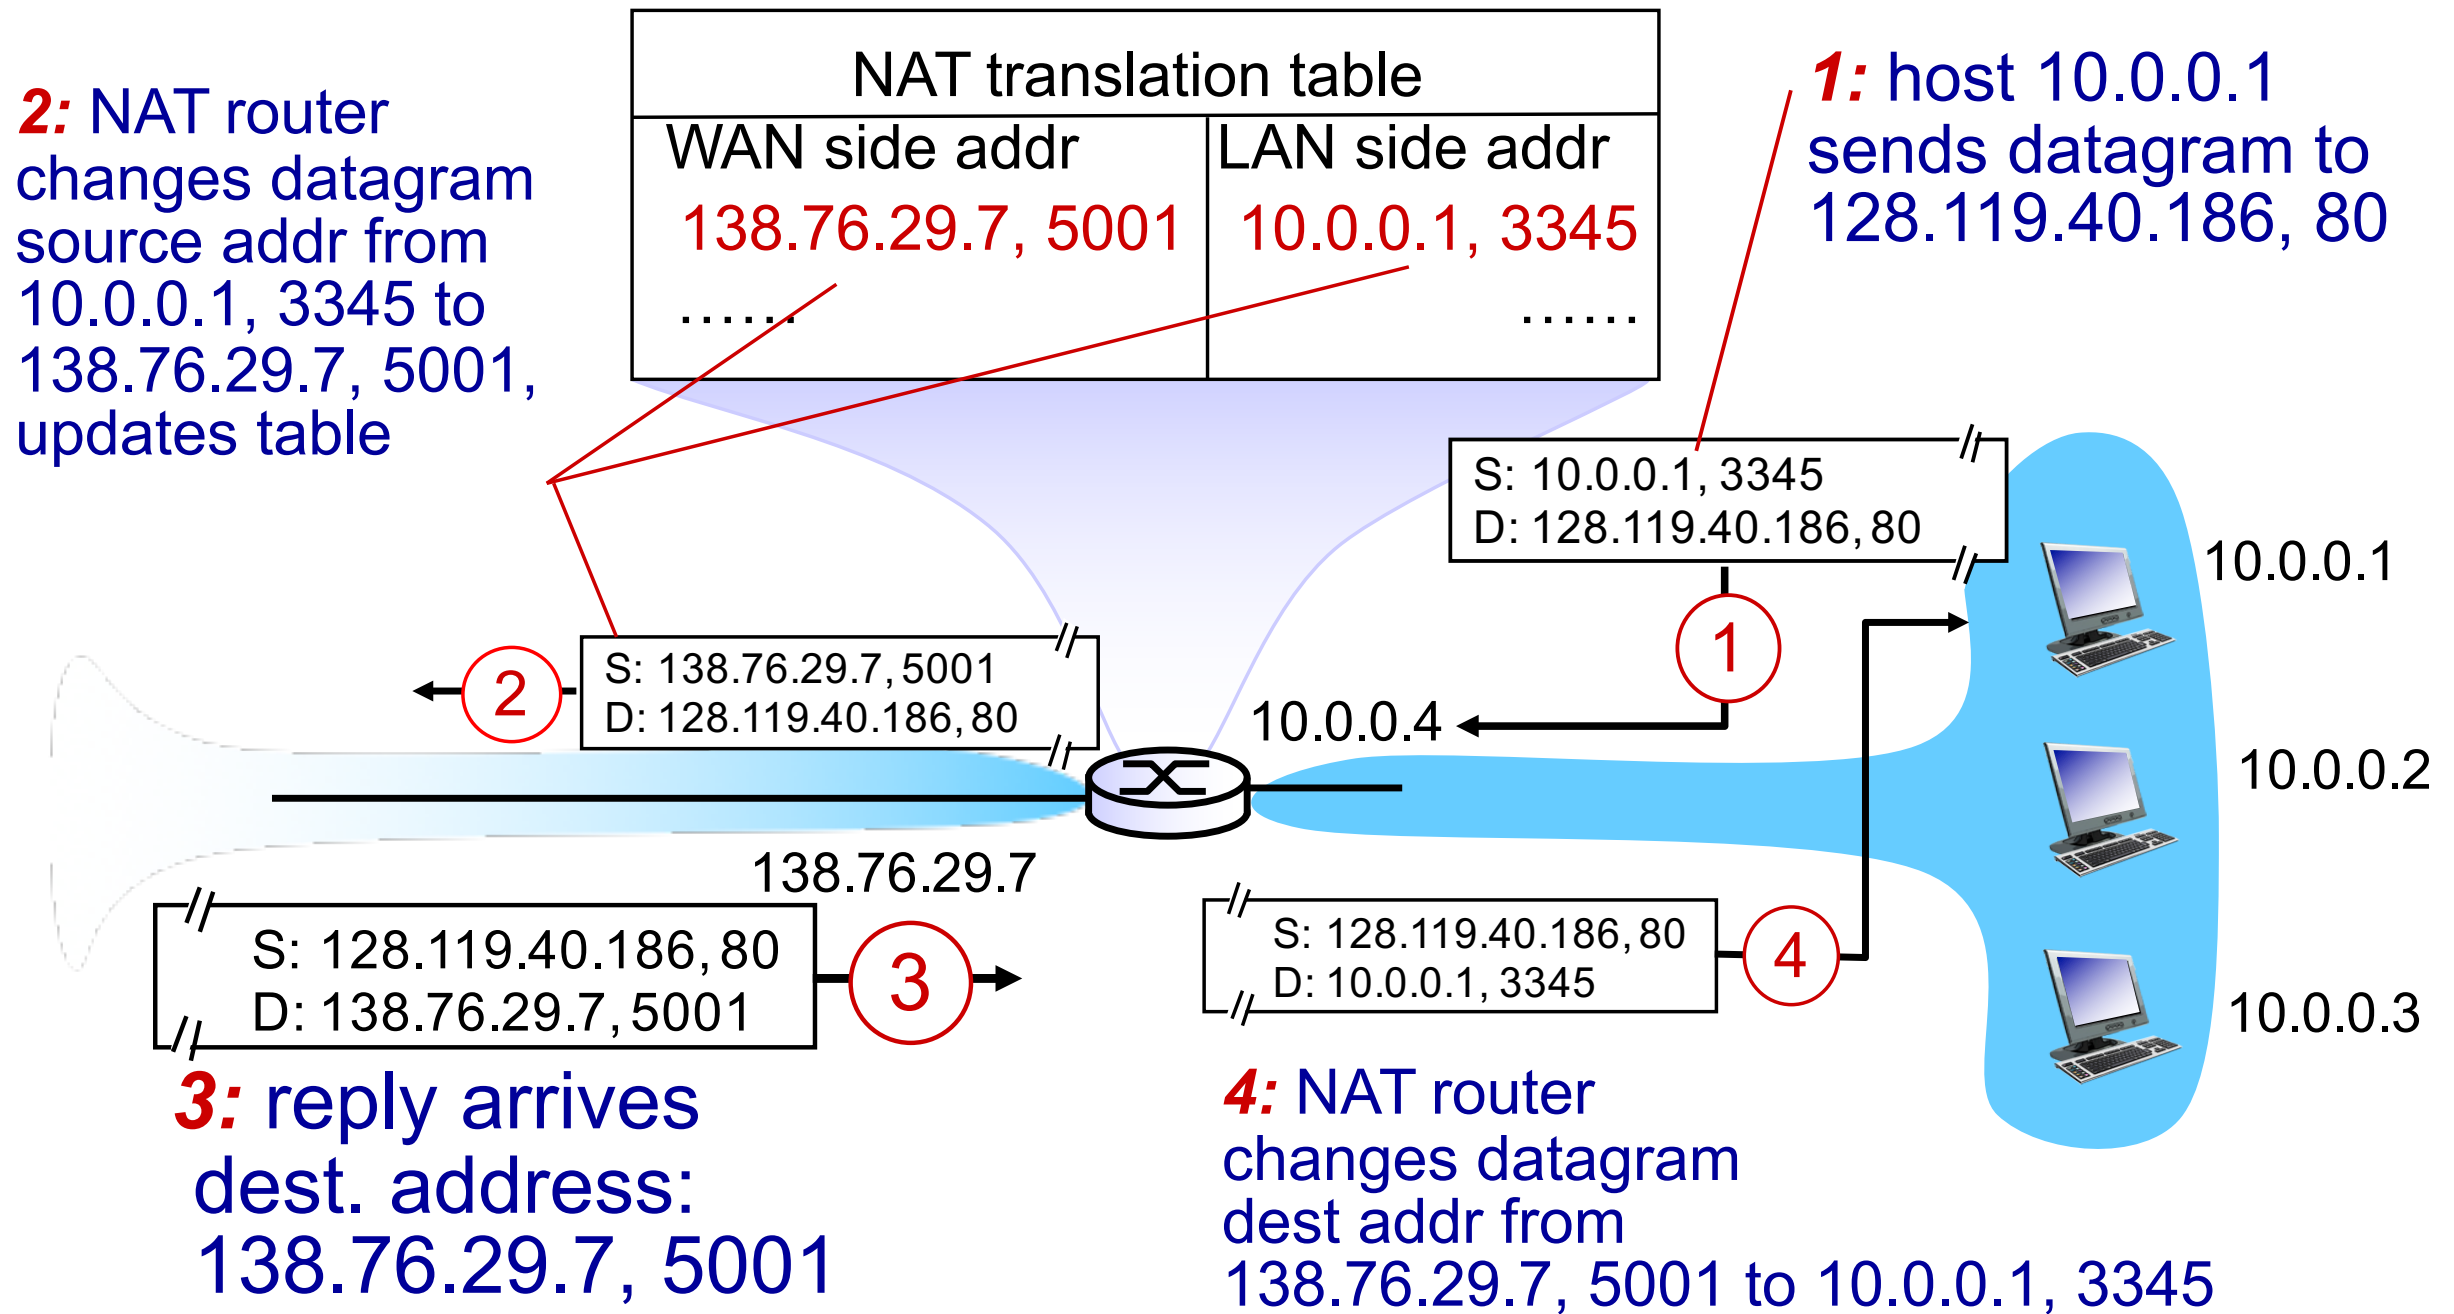
- range of addresses not needed from ISP: just one IP address for all devices
- can change addresses of devices in local network without notifying outside world
- can change ISP without changing addresses of devices in local network
- devices inside local net not explicitly addressable, visible by outside world (a security plus)

NAT: network address translation

implementation: NAT router must:

- *outgoing datagrams: replace* (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)
... remote clients/servers will respond using (NAT IP address, new port #) as destination addr
- *remember (in NAT translation table)* every (source IP address, port #) to (NAT IP address, new port #) translation pair
- *incoming datagrams: replace* (NAT IP address, new port #) in dest fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

NAT: network address translation



NAT: network address translation

- ❖ 16-bit port-number field:
 - 60,000 simultaneous connections with a single LAN-side address!
- ❖ NAT is controversial:
 - routers should only process up to layer 3
 - violates end-to-end argument
 - NAT possibility must be taken into account by app designers, e.g., P2P applications
 - address shortage should instead be solved by IPv6

IPv6: motivation

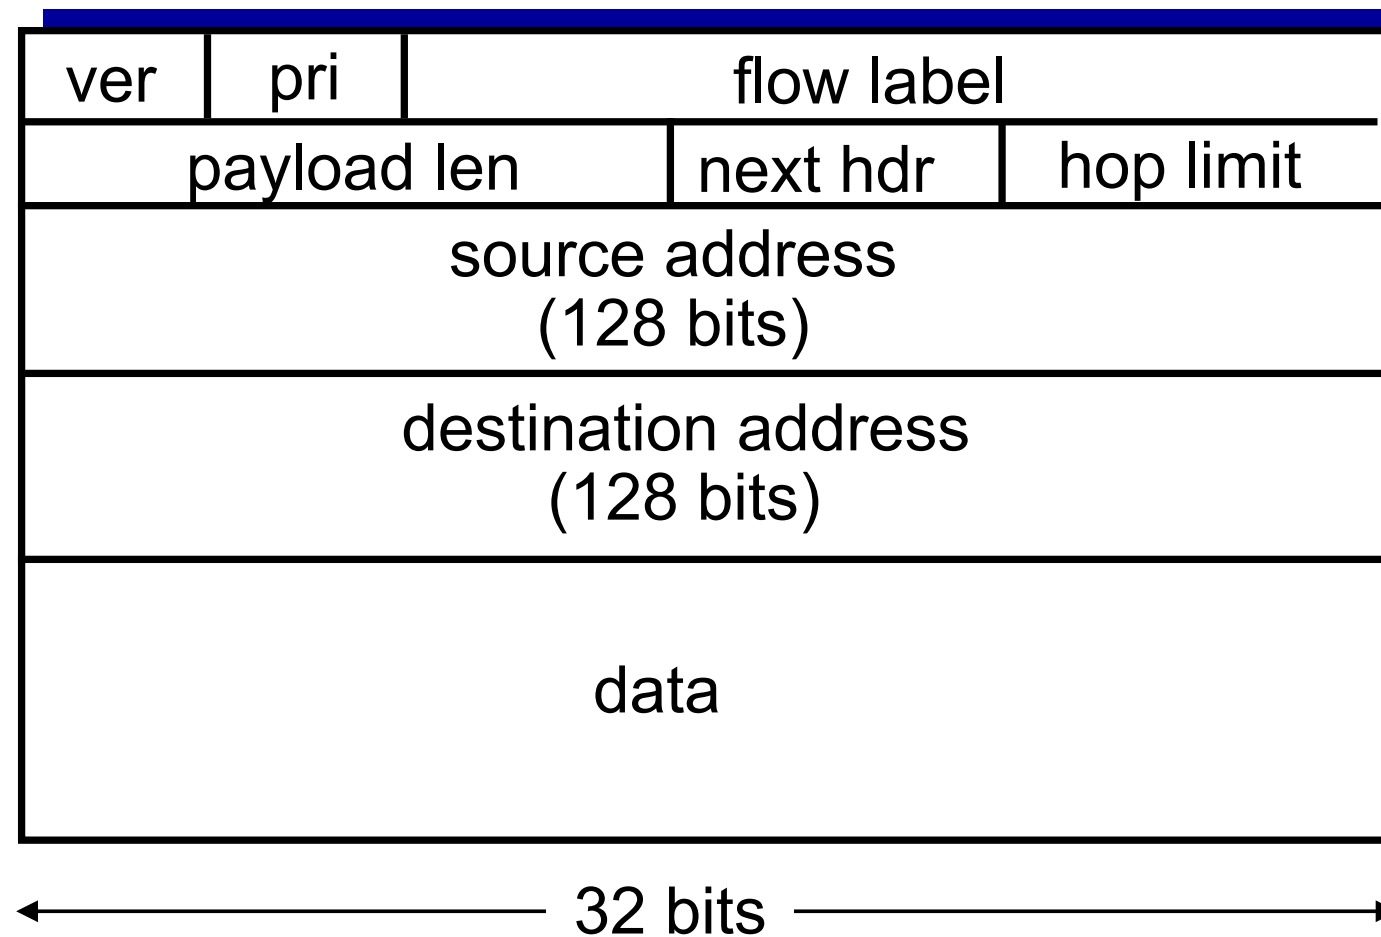
- ❖ *initial motivation*: 32-bit address space soon to be completely allocated.
- ❖ additional motivation:
 - header format helps speed processing/forwarding
 - header changes to facilitate QoS

IPv6 datagram format:

- fixed-length 40 byte header
- no fragmentation allowed

IPv6 datagram format

- priority*: identify priority among datagrams in flow
- flow Label*: identify datagrams in same “flow.”
(concept of “flow” not well defined).
- next header*: identify upper layer protocol for data



Network Interface Configuration

- Command for network interface configuration:
- `ifconfig [option] [interface]`
- If command not found, install the tool:

```
[root@localhost ~]# yum install net-tools
Loaded plugins: fastestmirror
Loading mirror speeds from cached hostfile
* base: mirror.lug.udel.edu
* extras: mirror.es.its.nyu.edu
* updates: mirrors.lga7.us.voxel.net
```

Network Interface Configuration

- Usage: check all network interfaces:
- `ifconfig -a`

```
[root@localhost ~]# ifconfig -a
enp0s3: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
        ether 08:00:27:12:ea:9a  txqueuelen 1000  (Ethernet)
        RX packets 0  bytes 0 (0.0 B)
        RX errors 0  dropped 0  overruns 0  frame 0
        TX packets 0  bytes 0 (0.0 B)
        TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0

enp0s8: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
        inet 192.168.1.31  netmask 255.255.255.0  broadcast 192.168.1.255
        inet6 fe80::ae7e:dca9:dfa6:fabe  prefixlen 64  scopeid 0x20<link>
        ether 08:00:27:34:2b:0c  txqueuelen 1000  (Ethernet)
        RX packets 10374  bytes 14649660 (13.9 MiB)
        RX errors 0  dropped 0  overruns 0  frame 0
        TX packets 1512  bytes 127449 (124.4 KiB)
        TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0

lo:  flags=73<UP,LOOPBACK,RUNNING>  mtu 65536
        inet 127.0.0.1  netmask 255.0.0.0
        inet6 ::1  prefixlen 128  scopeid 0x10<host>
        loop txqueuelen 1  (Local Loopback)
        RX packets 0  bytes 0 (0.0 B)
        RX errors 0  dropped 0  overruns 0  frame 0
        TX packets 0  bytes 0 (0.0 B)
        TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0
```

Network Interface Configuration

- Understanding output information:

```
[root@localhost ~]# ifconfig enp0s8
enp0s8: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
    inet 192.168.1.31  netmask 255.255.255.0  broadcast 192.168.1.255
    inet6 fe80::ae7e:dca9:dfa6:fabe  prefixlen 64  scopeid 0x20<link>
    ether 08:00:27:34:2b:0c  txqueuelen 1000  (Ethernet)
    RX packets 10494  bytes 14663846 (13.9 MiB)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 1566  bytes 134435 (131.2 KiB)
    TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0
```

- enp0s8: name of the device
- UP: This flag indicates that the kernel modules related to the Ethernet interface has been loaded.
- BROADCAST: Denotes that the Ethernet device supports broadcasting - a necessary characteristic to obtain IP address via DHCP.
- RUNNING: The interface is ready to accept data

Network Interface Configuration

- Understanding output information:

```
[root@localhost ~]# ifconfig enp0s8
enp0s8: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
    inet 192.168.1.31  netmask 255.255.255.0  broadcast 192.168.1.255
    inet6 fe80::ae7e:dca9:dfa6:fabe  prefixlen 64  scopeid 0x20<link>
    ether 08:00:27:34:2b:0c  txqueuelen 1000  (Ethernet)
    RX packets 10494  bytes 14663846 (13.9 MiB)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 1566  bytes 134435 (131.2 KiB)
    TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0
```

- MULTICAST: This indicates that the Ethernet interface supports multicasting.
- mtu: short form for Maximum Transmission Unit is the size of each packet received by the Ethernet card.
- inet: IPv4 address
- inet6: IPv6 address
- ether: MAC address
- txqueuelen: This denotes the length of the transmit queue of the device.

Network Interface Configuration

- Understanding output information:

```
[root@localhost ~]# ifconfig enp0s8
enp0s8: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
    inet 192.168.1.31  netmask 255.255.255.0  broadcast 192.168.1.255
    inet6 fe80::ae7e:dca9:dfa6:fabe  prefixlen 64  scopeid 0x20<link>
    ether 08:00:27:34:2b:0c  txqueuelen 1000  (Ethernet)
    RX packets 10494  bytes 14663846 (13.9 MiB)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 1566  bytes 134435 (131.2 KiB)
    TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0
```

- RX and TX: total number of packets or bytes received and transmitted respectively.

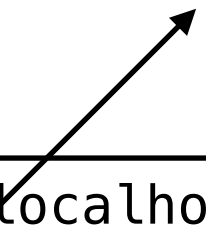
Network Interface Configuration

- Change IP address using ifconfig:
- *ifconfig <interface> <address> netmask <netmask> [up|down]*
- *ifconfig <interface> <address> </prefixlen> [up | down]*
-

```
[root@localhost ~]# ifconfig enp0s3 192.168.2.3/24
[root@localhost ~]# ifconfig enp0s3
enp0s3: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
    inet 192.168.2.3  netmask 255.255.255.0  broadcast 192.168.2.255
    inet6 fe80::a00:27ff:fe12:ea9a  prefixlen 64  scopeid 0x20<link>
    ether 08:00:27:12:ea:9a  txqueuelen 1000  (Ethernet)
    RX packets 17  bytes 3135 (3.0 KiB)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 542  bytes 33328 (32.5 KiB)
    TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0
```

Network Interface Configuration

- ifconfig:
 - changes network interface configurations temporarily
 - lose configurations after reboot
- Edit network configuration files:
 - Location: /etc/sysconfig/network-scripts/
 - Files start with "ifcfg-interfaceName" are configuration files for network interfaces



```
[root@localhost network-scripts]# ls
ifcfg-enp0s3  ifdown-isdn  ifdown-tunnel  ifup-isdn  ifup-Team
ifcfg-lo      ifdown-post  ifup           ifup-plip  ifup-TeamPort
ifdown       ifdown-ppp   ifup-aliases  ifup-plusb ifup-tunnel
ifdown-bnep  ifdown-routes ifup-bnep      ifup-post  ifup-wireless
ifdown-eth   ifdown-sit   ifup-eth       ifup-ppp   init.ipv6-global
ifdown-ipp   ifdown-Team  ifup-ipp       ifup-routes network-functions
ifdown-ipv6  ifdown-TeamPort ifup-ipv6      ifup-sit   network-functions-ipv6
```

Network Interface Configuration

- Understand the configuration file:
 - BOOTPROTO:
 - boot-time protocol used
 - none - No boot time protocol should be used
 - dhcp - The DHCP protocol should be used
 - static - Use static ip address
 - ONBOOT:
 - yes - This device should be activated at boot-time
 - no - This device should not be activated at boot-time

```
TYPE=Ethernet
PROXY_METHOD=none
BROWSER_ONLY=no
BOOTPROTO=dhcp
DEFROUTE=yes
IPV4_FAILURE_FATAL=no
IPV6INIT=yes
IPV6_AUTOCONF=yes
IPV6_DEFROUTE=yes
IPV6_FAILURE_FATAL=no
IPV6_ADDR_GEN_MODE=stable-privacy
NAME=enp0s3
UUID=dd33654a-838f-4664-abb4-38cf1b5980ac
DEVICE=enp0s3
ONBOOT=no
```


Network Interface Configuration

- Config a static IP interface:
 - IPADDR: IPv4 address
 - NETMASK: network mask

```
TYPE=Ethernet
PROXY_METHOD=none
BROWSER_ONLY=no
BOOTPROTO=static
IPADDR=192.168.2.3
NETMASK=255.255.255.0
NAME=enp0s3
UUID=dd33654a-838f-4664-
abb4-38cf1b5980ac
DEVICE=enp0s3
ONBOOT=yes
```

~

Network Interface Configuration

- Take a network interface down:

```
[root@localhost network-scripts]# ifdown enp0s3  
Device 'enp0s3' successfully disconnected.
```

- Bring a network interface up:

```
[root@localhost network-scripts]# ifup enp0s3  
Connection successfully activated (D-Bus active path: /org/freedesktop/  
NetworkManager/ActiveConnection/3)
```

Hostname

- Change your hostname by command:
 - Temporary

```
[root@localhost ~]# hostname  
localhost.localdomain  
[root@localhost ~]# hostname myhost.localdomain  
[root@localhost ~]# hostname  
myhost.localdomain
```

- Change your hostname by editing configuration file:
- Location: /etc/hostname

```
[root@localhost ~]# echo myhost.localhost > /etc/hostname
```

Hosts files

- /etc/hosts: file retains a mapping of host names and their ip addresses

-

```
[root@localhost ~]# cat /etc/hosts
127.0.0.1    localhost localhost.localdomain localhost4 localhost4.localdomain4
::1         localhost localhost.localdomain localhost6 localhost6.localdomain6
```

Network Tools

- ping: used to test connections between hosts
- traceroute: used to test connections between hosts and routers
- netstat: Print network connections, routing tables, interface statistics, masquerade connections, and multicast memberships

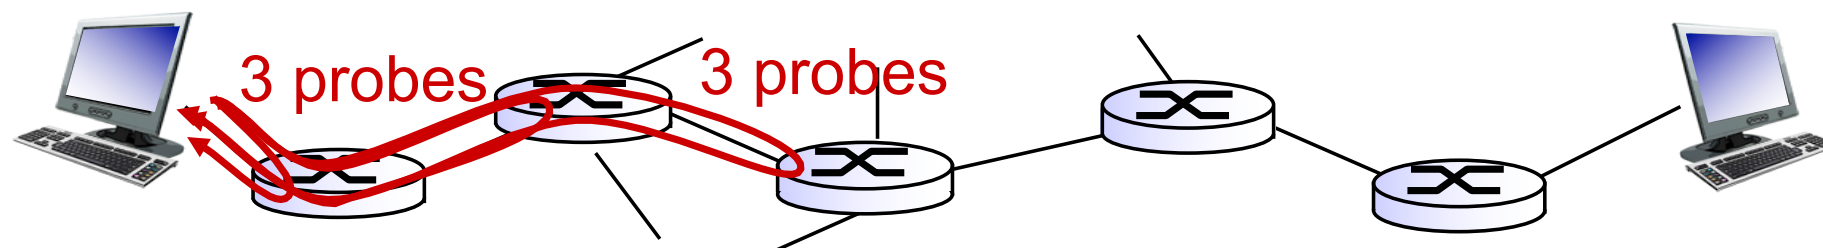
Traceroute and ICMP

- ❖ source sends series of UDP segments to dest
 - first set has TTL = 1
 - second set has TTL=2, etc.
 - unlikely port number
- ❖ when n th set of datagrams arrives to n th router:
 - router discards datagrams
 - and sends source ICMP messages (type 11, code 0)
 - ICMP messages includes name of router & IP address

- ❖ when ICMP messages arrives, source records RTTs

stopping criteria:

- ❖ UDP segment eventually arrives at destination host
- ❖ destination returns ICMP “port unreachable” message (type 3, code 3)
- ❖ source stops



netstat

```
[root@localhost ~]# netstat -tunpa
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address           Foreign Address         State       PID/Program name
tcp        0      0 0.0.0.0:22              0.0.0.0:*               LISTEN      964/sshd
tcp        0      0 127.0.0.1:25            0.0.0.0:*               LISTEN      1066/master
tcp        0      0 192.168.1.31:22         192.168.1.8:61706      ESTABLISHED 1257/sshd: root@pts
tcp6       0      0 :::22                  :::*                   LISTEN      964/sshd
tcp6       0      0 :::1:25                 :::*                   LISTEN      1066/master
udp        0      0 0.0.0.0:68              0.0.0.0:*               783/dhclient
udp        0      0 0.0.0.0:4317            0.0.0.0:*               783/dhclient
udp        0      0 127.0.0.1:323           0.0.0.0:*               644/chronyd
udp6       0      0 :::13571                :::*                   783/dhclient
udp6       0      0 :::1:323                :::*                   644/chronyd
```

- Options:

- a: Show both listening and non-listening sockets.
- t: TCP
- u: UDP
- n: Show numerical addresses instead of trying to determine symbolic host, port or user names.
- p: Show the PID and name of the program to which each socket belongs.